

Getting Started with HBase: The Hadoop Database

INTRODUCING HBASE

Overview

Understand the need for HBase in a distributed environment

Understand the differences between HBase and an RDBMS

Install and set up HBase

Software for Distributed Computing

How Much Data Do Organizations Deal With?

The Google logo is centered within a solid green square. The word "Google" is written in its characteristic white, multi-colored sans-serif font.

Google

The Google logo is displayed in white text on a solid green rectangular background.

Google

Current storage = 15 exabytes


Processed per day = 100 petabytes

Number of pages indexed = 60 trillion

Unique search users per month > 1 billion



Searches per second = 2.3 million

Single Coordinating Software





**Google realized that
running web search
algorithms on distributed
systems required special
software**

Single Coordinating Software





1: Store millions of records on multiple machines

Single Coordinating Software




**2: Run processes on all these machines
to crunch data**

Single Coordinating Software



**3: Handle fault tolerance and recovery
when nodes crash**

Single Coordinating Software



Google File System

To solve
distributed
storage

MapReduce

To solve
distributed
computing

Single Coordinating Software

Google File System

MapReduce

**Apache developed
open source versions
of these technologies**

Single Coordinating Software

Google File System




HDFS

MapReduce



MapReduce

Hadoop



**A file system to manage
the storage of data**

**A framework to process data
across multiple servers**

Hadoop is a big data processing
framework

Hadoop is **not** a database!


The Importance of Databases

What Kind of Data Do Organizations Store?



Order Management

An e-commerce site stores order information




Payroll

A company stores employee payroll details



Accounts

A bank stores account and transaction information



Requirements of a Database

Structured: Rows and columns

Random access: Update one row at a time

Low latency: Very fast read/write/update operations

ACID compliant: Ensure data integrity

What Are ACID Properties?

Atomicity

Consistency

Isolation


Durability

Atomicity

Transactions on a
database should
be **all-or-nothing**

Atomicity

Transferring Money




**Both withdrawal and deposit
should occur or none at all!**

Consistency

Database updates
should not violate
any **constraints**

Consistency

Enrolling Students



**Every student should
have a unique student id**

Isolation

**Concurrent operations
on the database should
appear as though they
were applied in some
sequence**

Isolation

Granularity of Updates



Can employee address be updated at the same time as employee salary?

Durability

Once changes have
been made to the
data they are
permanent

Durability

Safety of Data



**In case of power
loss, crashes, errors**

The Limitations of Hadoop

Unfortunately, Hadoop makes a
very poor database

Limitations of Hadoop



Unstructured data



No random access



High latency



Not ACID compliant

Limitations of Hadoop

Data in HDFS has no schema,
no rows and columns, no tables

Text files

Log files

Audio files

Video files



Unstructured data

Limitations of Hadoop

**Basic structure exists
for some file types**

CSV files

XML files

JSON files

**Hadoop enforces no
constraints on these**



Unstructured data

Limitations of Hadoop

Cannot create, access
and modify individual
records in a file

MapReduce parses
entire files to extract
information



No random access

Limitations of Hadoop

Not suited for real-time processing where a user waits for data to be retrieved

Batch processing with long running jobs



High latency


Limitations of Hadoop

HDFS is a file storage system and provides no guarantees for data integrity




Not ACID compliant

How Did Google Solve This for Search?



**Google published a paper
on Bigtable a distributed
storage system for
structured data**

How Did Google Solve This for Search?



HBase is a distributed database management system which runs on top of Hadoop

HBase

The diagram illustrates the HBase architecture. At the top is a large gray rectangle labeled 'HBase'. Below it is a blue-bordered container. Inside this container are two gray rectangles: 'HDFS' on the left and 'MapReduce' on the right. This indicates that HBase interacts with both HDFS and MapReduce.

HBase

HDFS

MapReduce



HBase

Distributed: Stores data in HDFS

Scalable: Capacity directly proportional to number of nodes in the cluster

Fault tolerant: Piggybacks on Hadoop



HBase

Structured: A loose data structure

Low latency: Real-time access using row based indices called row keys

Random access: Row keys allow access updates to one record

Somewhat ACID compliant: Some transactions will have ACID properties




HBase

**Batch processing
using MapReduce**

**Real-time processing
using row keys**

HBase vs. Relational Databases

Properties of HBase



Columnar store



Denormalized storage



Only CRUD operations



ACID at the row level




Columnar store

A Notification Service

Id	To	Type	Content
1	mike	offer	Offer on mobiles
2	john	sale	Redmi sale
3	jill	order	Order delivered
4	megan	sale	Clothes sale

**Layout of a traditional
relational database**




Columnar store

A Notification Service

Id	To	Type	Content
1	mike	offer	Offer on mobiles
2	john	sale	Redmi sale
3	jill	order	Order delivered
4	megan	sale	Clothes sale

Layout of a traditional relational database



Columnar store

A Notification Service

Id	To	Type	Content
1	mike	offer	Offer on mobiles
2	john	sale	Redmi sale
3	jill	order	Order delivered
4	megan	sale	Clothes sale

Row = 3
Column = To

Columnar Store

Id	To	Type	Content
1	mike	offer	Offer on mobiles
2	john	sale	Redmi sale
3	jill	order	Order delivered
4	megan	sale	Clothes sale



Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on mobiles
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	megan
4	Type	sale
4	Content	Clothes sale

Columnar Store

Id	To	Type	Content
1	mike	offer	Offer on mobiles
2	john	sale	Redmi sale
3	jill	order	Order delivered
4	megan	sale	Clothes sale



Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on mobiles
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	megan
4	Type	sale
4	Content	Clothes sale

Columnar Store

Id	To	Type	Content
1	mike	offer	Offer on mobiles
2	john	sale	Redmi sale
3	jill	order	Order delivered
4	megan	sale	Clothes sale




Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on mobiles
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	megan
4	Type	sale
4	Content	Clothes sale

Columnar Store

Id	To	Type	Content
1	mike	offer	Offer on mobiles
2	john	sale	Redmi sale
3	jill	order	Order delivered
4	megan	sale	Clothes sale



Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on mobiles
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	megan
4	Type	sale
4	Content	Clothes sale



Columnar store

Advantages of a Columnar Store

Sparse tables: No wastage of space when storing sparse data

Dynamic attributes: Update attributes dynamically without changing storage structure

Sparse Tables

Id	To	Type	Content	Expiry
1	mike	offer	Offer on mobiles	2345689070
2	john	sale	Redmi sale	
3	jill	order	Order delivered	
4	megan	sale	Clothes sale	2456123989

**Sale and offer notifications
may have an expiry time**

Sparse Tables

Id	To	Type	Content	Expiry	Order Status
1	mike	offer	Offer on mobiles	2345689070	
2	john	sale	Redmi sale		
3	jill	order	Order delivered		Delivered
4	megan	sale	Clothes sale	2456123989	

**Order related notifications
may have an order status**

Sparse Tables

Id	To	Type	Content	Expiry	Order Status
1	mike	offer	Offer on mobiles	2345689070	
2	john	sale	Redmi sale		
3	jill	order	Order delivered		Delivered
4	megan	sale	Clothes sale	2456123989	

**In a traditional database this results
in a change in **database structure****

Sparse Tables


Id	To	Type	Content	Expiry	Order Status
1	mike	offer	Offer on mobiles	2345689070	
2	john	sale	Redmi sale		
3	jill	order	Order delivered		Delivered
4	megan	sale	Clothes sale	2456123989	

**And empty cells when data is
not applicable to certain rows**

Sparse Tables

Id	To	Type	Content	Expiry	Order Status
1	mike	offer	Offer on mobiles	2345689070	
2	john	sale	Redmi sale		
3	jill	order	Order delivered		Delivered
4	megan	sale	Clothes sale	2456123989	


These cells still occupy space!



Columnar store

Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on
1	Expiry	2345689070
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	megan
4	Type	sale
4	Content	Clothes sale
4	Expiry	2456123989


Dynamically add new attributes as rows in this table



Columnar store

Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on
1	Expiry	2345689070
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	megan
4	Type	sale
4	Content	Clothes sale
4	Expiry	2456123989

**No wastage of space
with empty cells!**




Columnar store

Note that this is not the exact layout of how data is stored in HBase

It is a **general structure** of how columnar stores are constructed

Properties of HBase



Columnar store



Denormalized storage



Only CRUD operations



ACID at the row level



Denormalized storage

**Traditional databases
use normalized forms
of database design to
minimize redundancy**



Denormalized storage

Minimize Redundancy

Employee Details

Employee Subordinates

Employee Address



Denormalized storage

Employee Details

Id	Name	Function	Grade
1	Emily	Finance	6

Employee Subordinates

Id	Subordinate Id
1	2
1	3

Employee Address

Id	City	Zip Code
1	Palo Alto	94305
2	Seattle	98101



Denormalized storage

Employee Details

Id	Name	Function	Grade
1	Emily	Finance	6
2	John	Finance	3
3	Ben	Finance	4

**All employee details
in one table**



Denormalized storage

Employee Subordinates

id	Subordinate id
1	2
1	3

**Employees referenced only
by ids everywhere else**



Denormalized storage

Employee Address

Id	City	Zip Code
1	Palo Alto	94305
2	Seattle	98101

Data is made more granular by splitting it across multiple tables



Denormalized storage

Id	Name	Function	Grade
1	Emily	Finance	6

Id	Subordinate Id
1	2
1	3

Id	City	Zip Code
1	Palo Alto	94305
2	Seattle	98101

Normalization

Normalization

Optimizes storage



Denormalized storage

**But storage is cheap in
a distributed system!**



Denormalized storage

**But storage is cheap in
a distributed system!**

**Optimize
number of disk
seeks**

Denormalized Storage

Id	Name	Function	Grade
1	Emily	Finance	6
2	John	Finance	3
3	Ben	Finance	4

Id	Subordinate Id
1	2
1	3



Id	Name	Function	Grade	Subordinates
1	Emily	Finance	6	<ARRAY>
2	John	Finance	3	
3	Ben	Finance	4	

Denormalized Storage

Id	Name	Function	Grade
1	Emily	Finance	6
2	John	Finance	3
3	Ben	Finance	4

Id	City	Zip Code
1	Palo Alto	94305
2	Seattle	98101



Id	Name	Function	Grade	Subordinates	Address
1	Emily	Finance	6	<ARRAY>	<STRUCT>
2	John	Finance	3		
3	Ben	Finance	4		

Denormalized Storage

Id	Name	Function	Grade	Subordinates	Address
1	Emily	Finance	6	<ARRAY>	<STRUCT>
2	John	Finance	3		
3	Ben	Finance	4		


Store everything related to an employee in the same table

Denormalized Storage

Id	Name	Function	Grade	Subordinates	Address
1	Emily	Finance	6	<ARRAY>	<STRUCT>
2	John	Finance	3		
3	Ben	Finance	4		

Read a single record to get all details about an employee in one read operation

Properties of HBase



Columnar store



Denormalized storage



Only CRUD operations



ACID at the row level



Only CRUD operations

Traditional Databases and SQL

Joins: Combining information across tables using keys

Group By: Grouping and aggregating data for the groups

Order By: Sorting rows by a certain column



Only CRUD operations

**HBase does
not support
SQL**

NoSQL



Only CRUD operations

**Only a limited set of operations
are allowed in HBase**

Create

Read

Uppdate

Delete

CRUD



Only CRUD operations

**No operations involving
multiple tables**

No indexes on tables

No constraints




Only CRUD operations

Id	Name	Function	Grade	Subordinates	Address

This is why all details
need to be **self**
contained in one row

Properties of HBase



Columnar store



Denormalized storage



Only CRUD operations



ACID at the row level



ACID at the row level

**Updates to a single
row are atomic**

**All columns in a
row are updated
or none are**



ACID at the row level

Updates to multiple
rows are **not** atomic

Even if the update is
on the same column
in multiple rows

Traditional RDBMS vs. HBase

Traditional RDBMS

Data arranged in rows and columns

Supports SQL

Complex queries such as grouping, aggregates, joins etc

Normalized storage to minimize redundancy and optimize space

ACID compliant

HBase

Data arranged in a column-wise manner

NoSQL database

Only basic operations such as create, read, update and delete

Denormalized storage to minimize disk seeks

ACID compliant at the row level

Demo

Install and set up HBase in pseudo-distributed mode

Summary

Understood the need for a distributed database system like HBase

Know how HBase differs from a traditional RDBMS

Installed and set up HBase on your local machine