

Integrating MapReduce with HBase

Overview

Understand the Hadoop MapReduce framework

Use MapReduce with HBase for complex SQL-like operations

Set up and run MapReduce on an HBase table

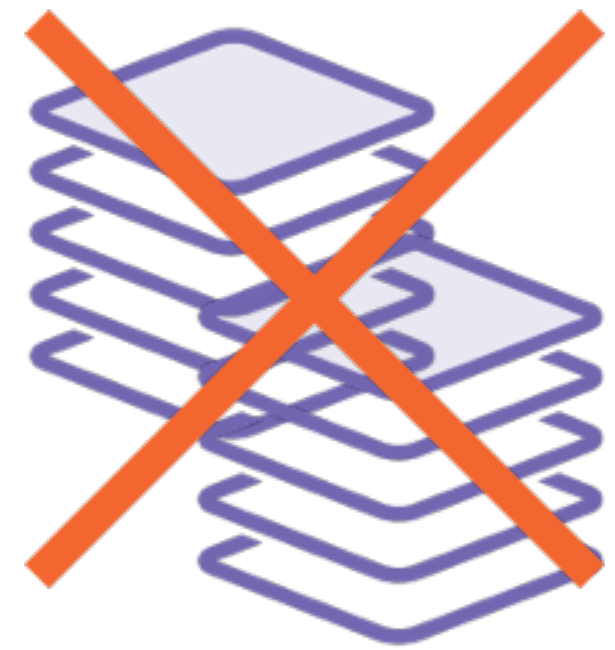
Complex and Multi-table Operations in HBase



**Grouping and
Sorting by column**



**Aggregation of
column values**



**Joins across
multiple tables**

Complex and Multi-table Operations in HBase

**There is a way
around this though...**

HBase allows **programmatic**
access to data using **MapReduce**

MapReduce

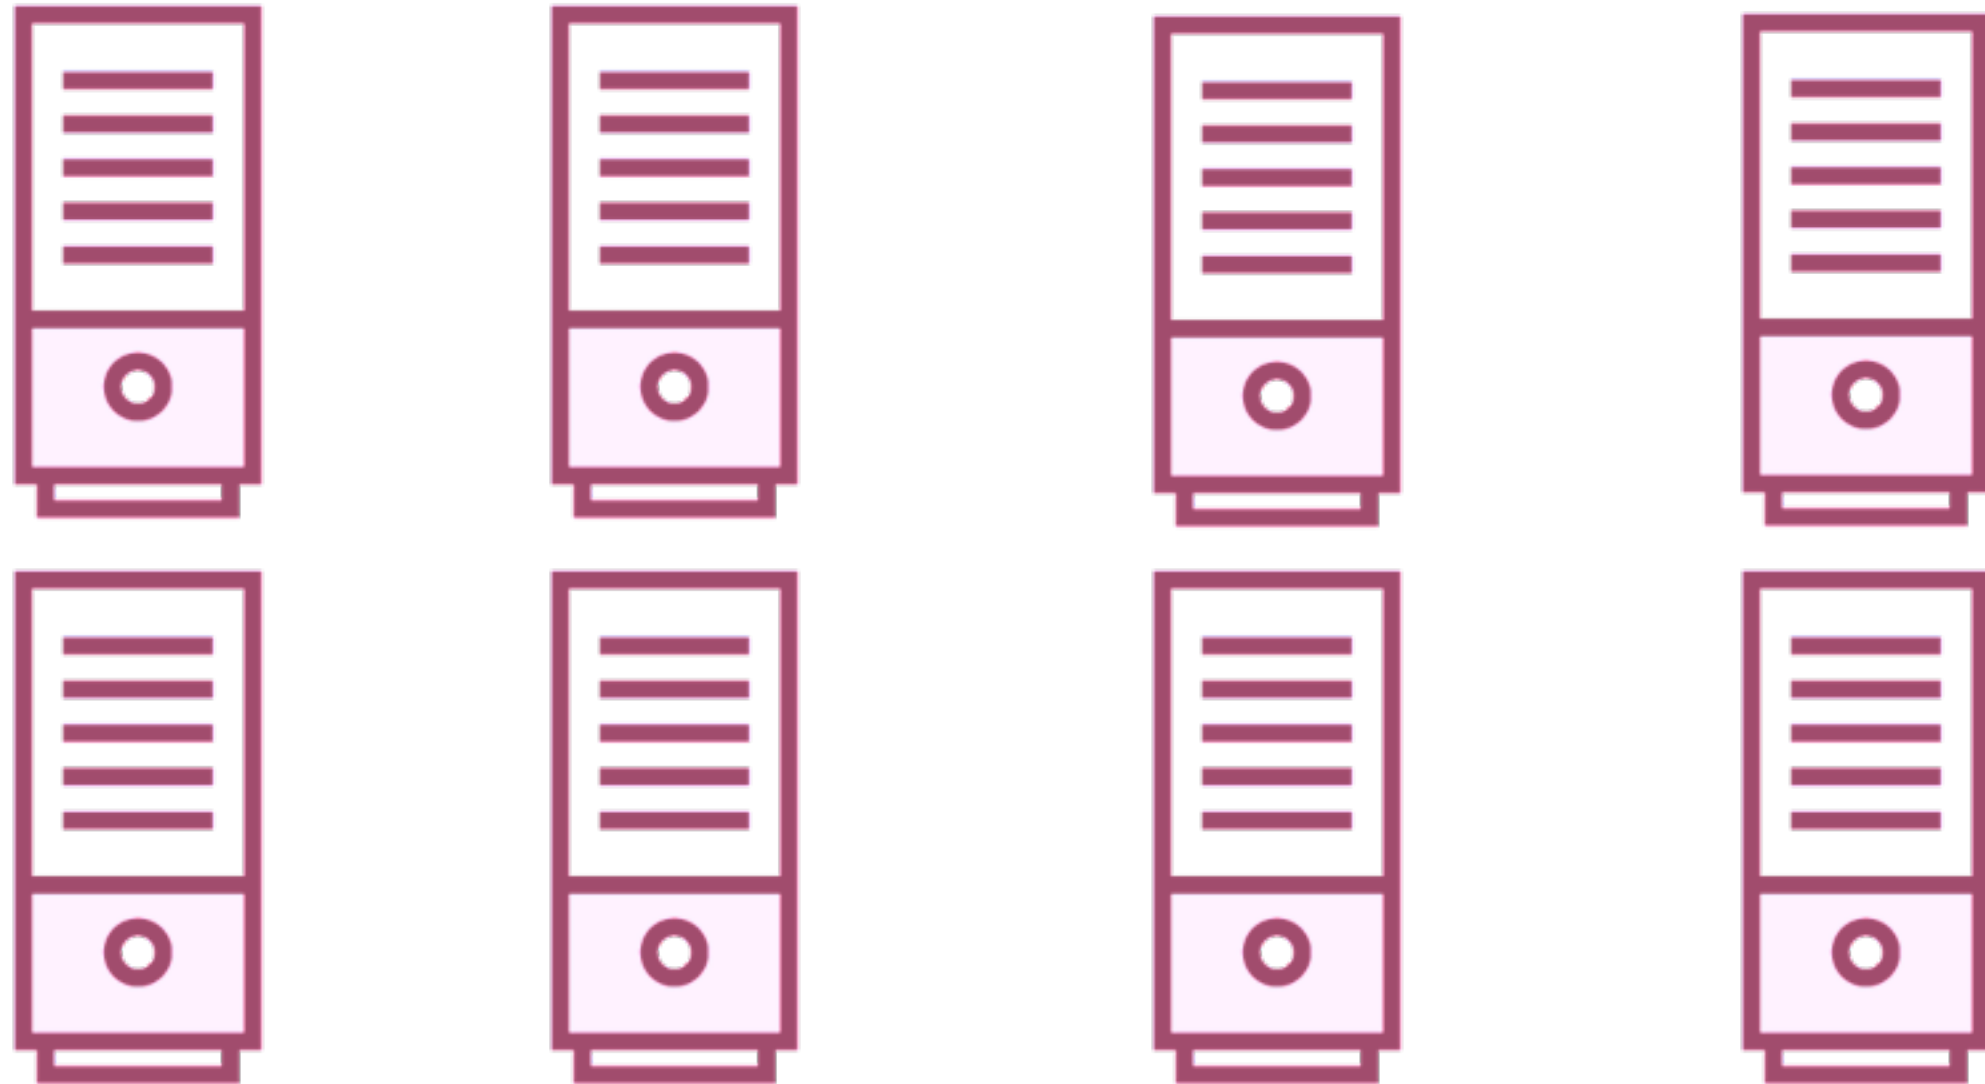
**Parallel programming model
usually implemented in Java**

**Full control over
how we process data**

MapReduce

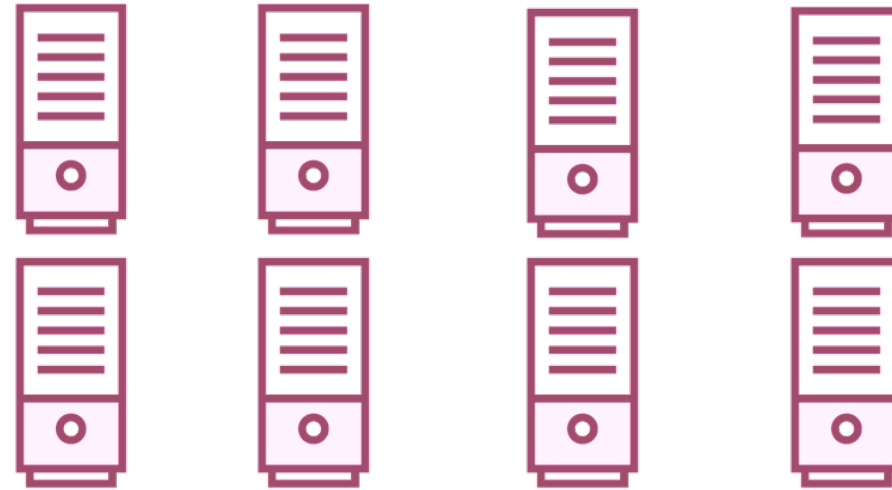
Processing huge amounts of data

MapReduce



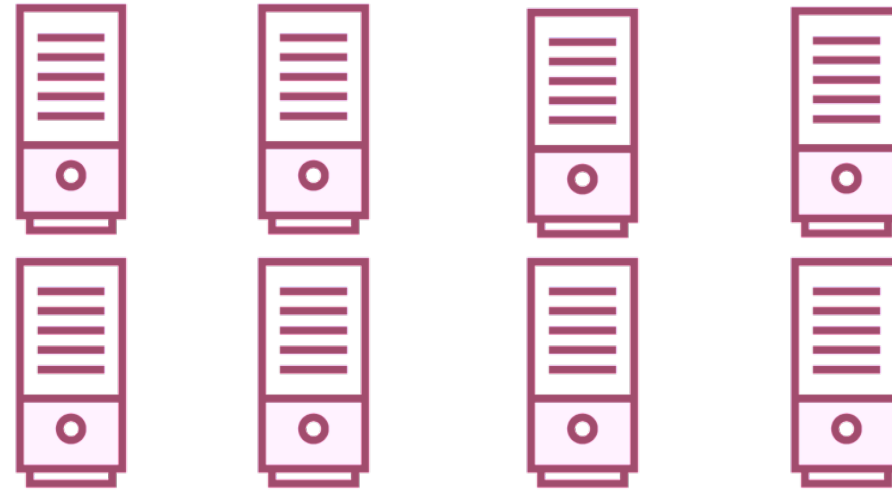
A distributed system

MapReduce



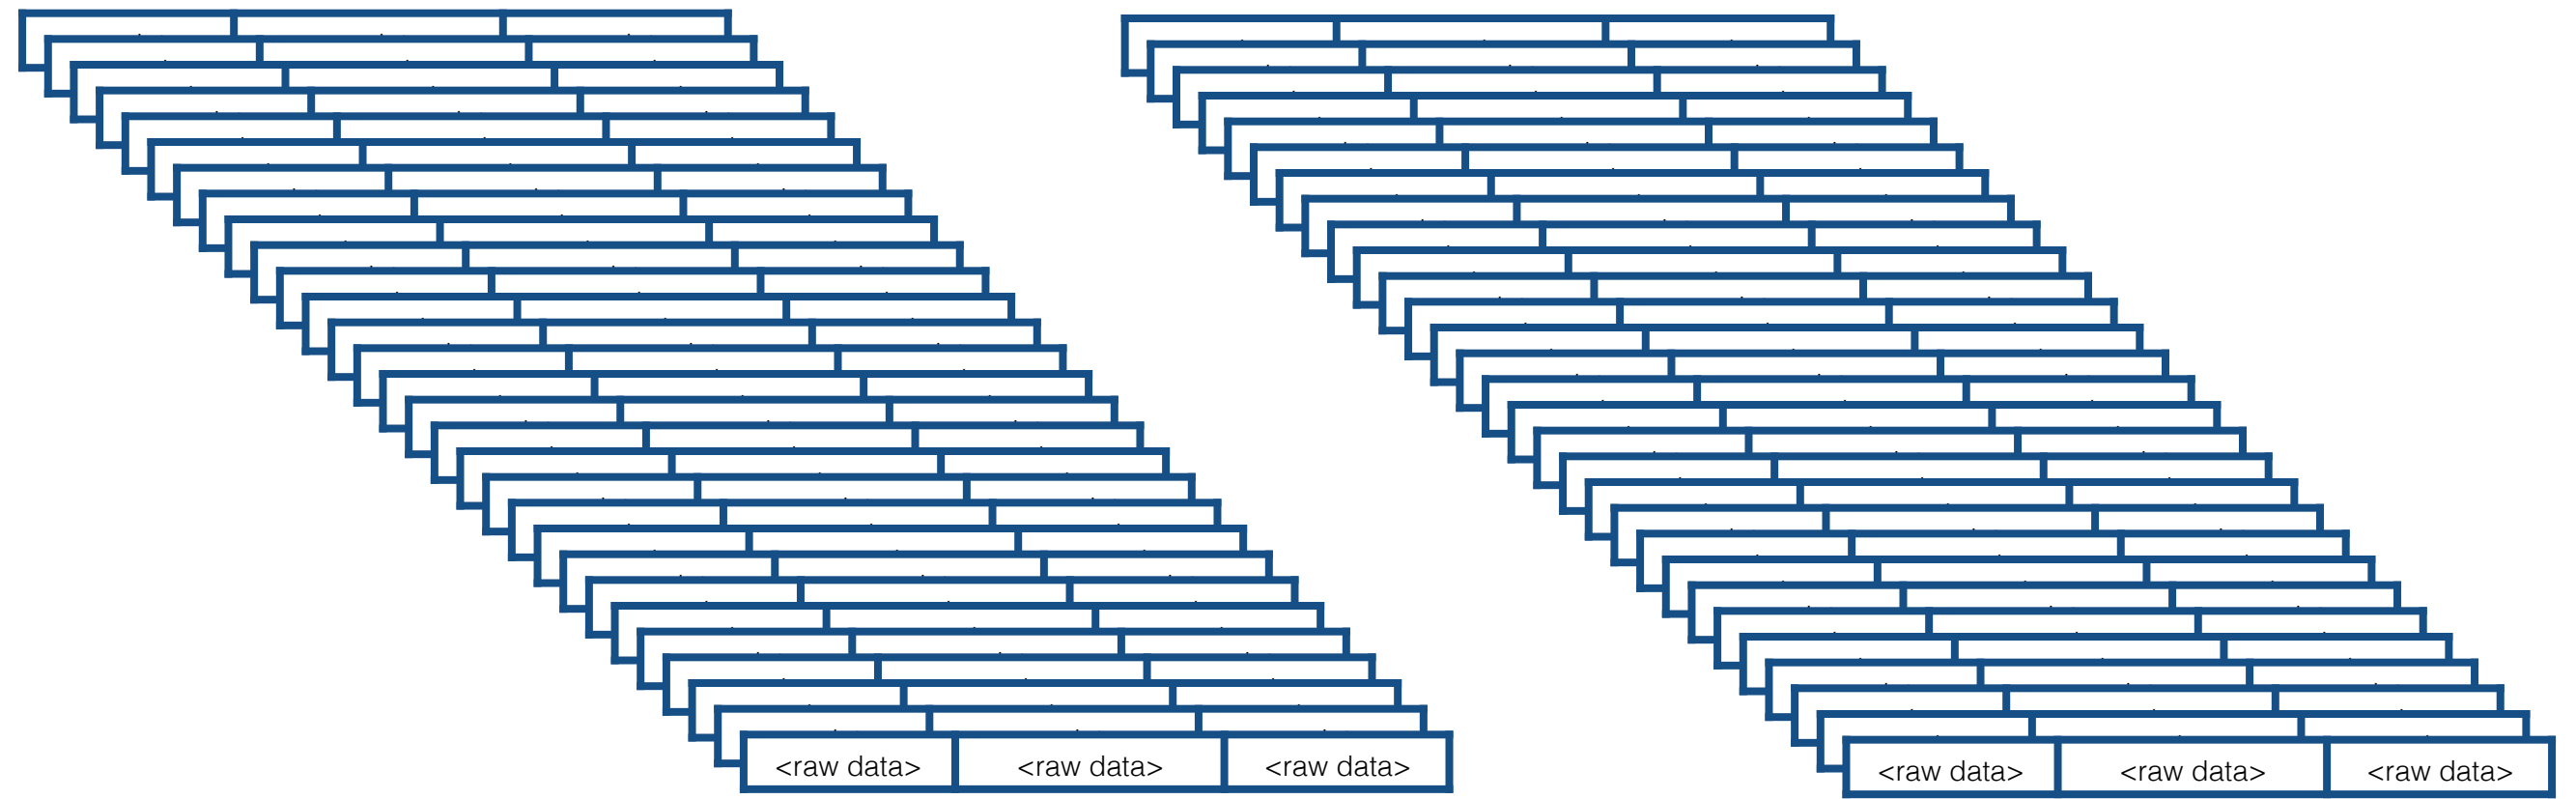
**MapReduce is a programming
paradigm**

MapReduce




**Takes advantage of the inherent
parallelism in data processing**

MapReduce



Modern systems generate millions of records of raw data

MapReduce


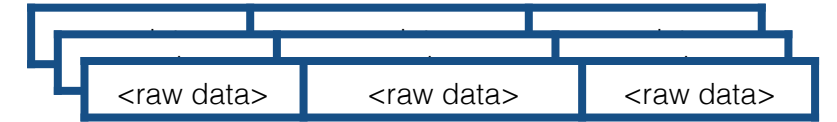
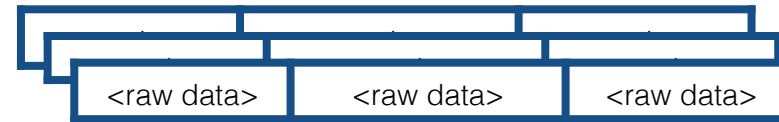
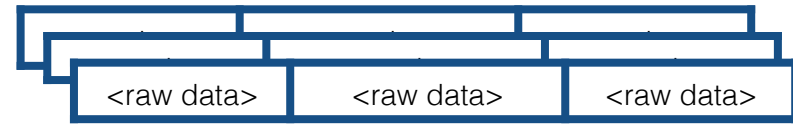


A task of this scale is processed in
two stages

map

reduce

map



reduce



<raw data>	<raw data>	<raw data>
<raw data>	<raw data>	<raw data>
<raw data>	<raw data>	<raw data>
<raw data>	<raw data>	<raw data>



MapReduce

map reduce

The programmer defines
these 2 functions

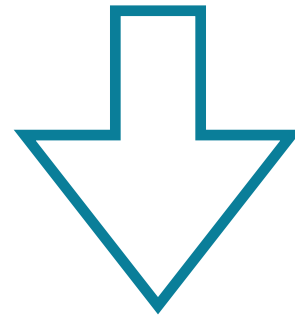
Hadoop does the rest -
behind the scenes

map

**An operation performed
in parallel, on small
portions of the dataset**

map

One Record

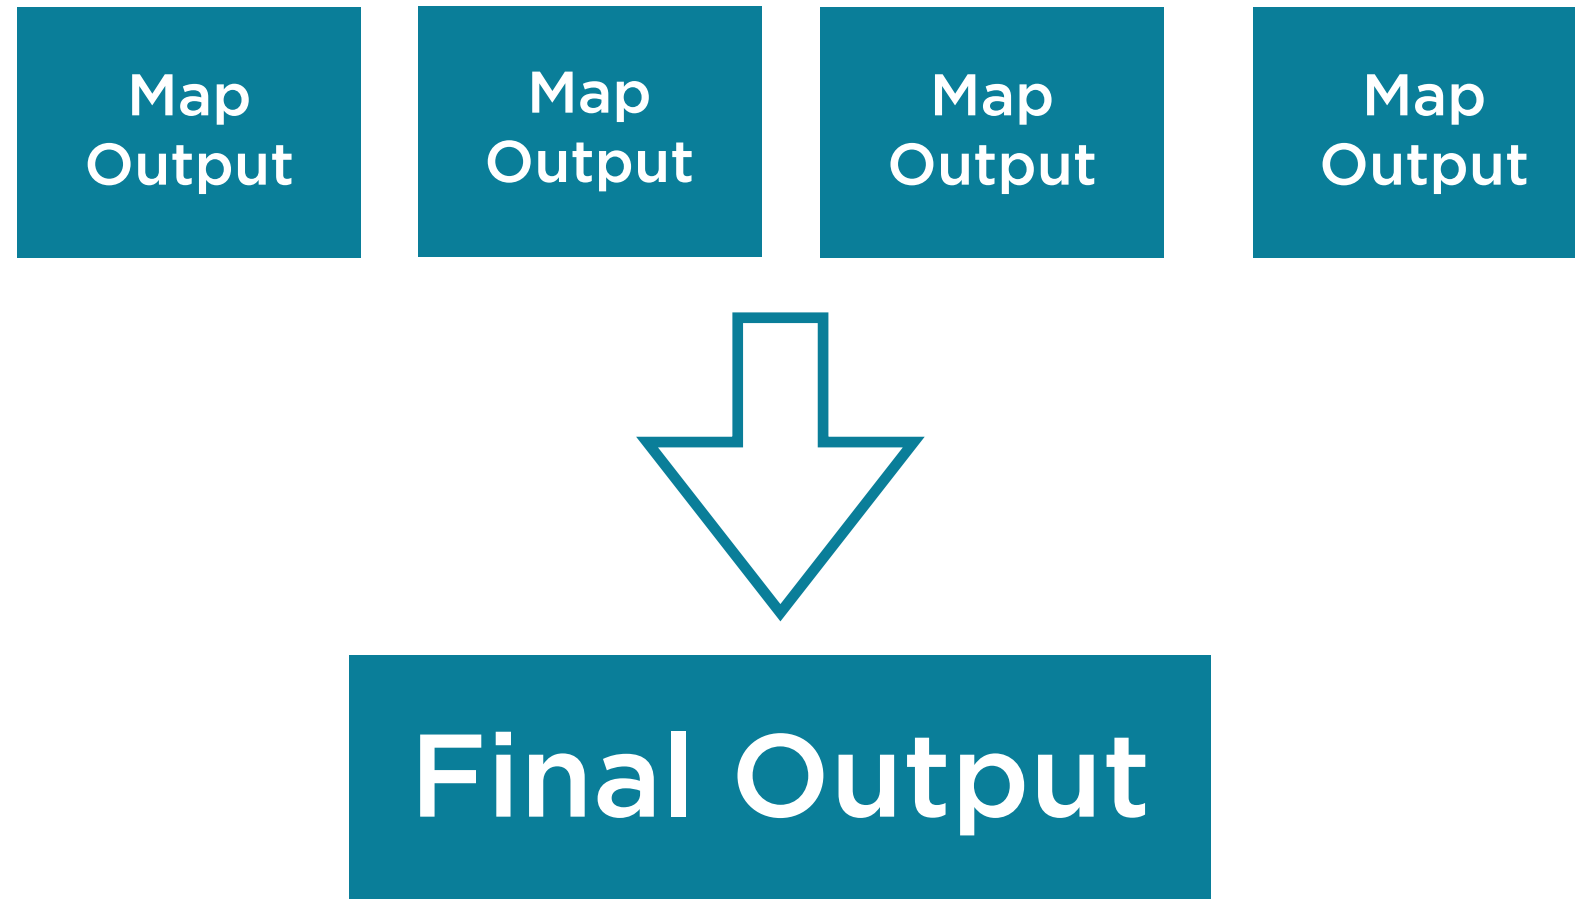


Key-Value Output

reduce

**An operation to
combine the results of
the map step**

reduce



map A step that can be
performed in parallel

reduce A step to combine the
intermediate results

Key Insight Behind MapReduce



Many data processing tasks can be expressed in this form

Key-Value Representation of an HBase Table

Key

Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on mobiles
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	mike
4	Type	sale
4	Content	Clothes sale

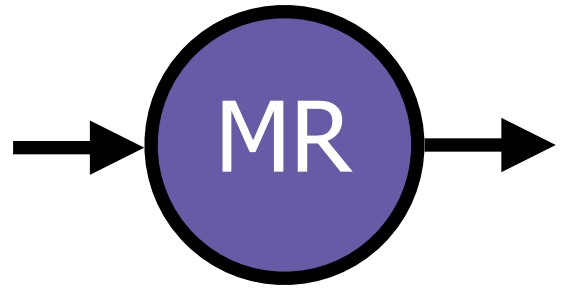
Key-Value Representation of an HBase Table

Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on mobiles
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	mike
4	Type	sale
4	Content	Clothes sale

Value

Message Count per User

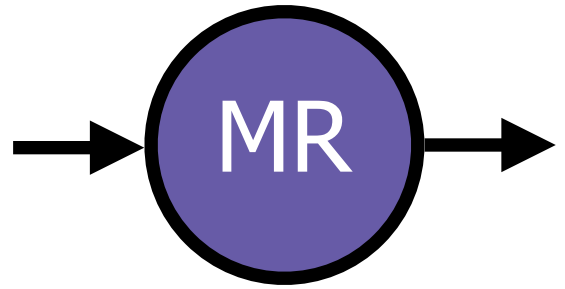
Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on mobiles
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	mike
4	Type	sale
4	Content	Clothes sale



To	Count
mike	2
john	1
jill	1

Message Count per User

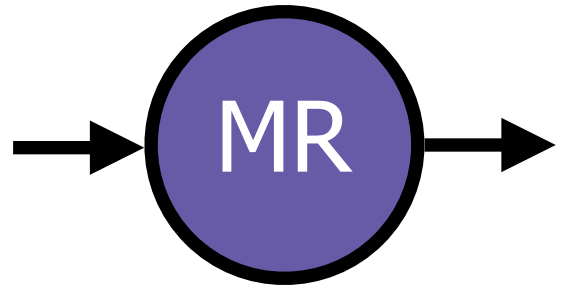
Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on mobiles
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	mike
4	Type	sale
4	Content	Clothes sale



To	Count
mike	2
john	1
jill	1

Message Count per User


Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on mobiles
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	mike
4	Type	sale
4	Content	Clothes sale



To	Count
mike	2
john	1
jill	1

Notification Type Count

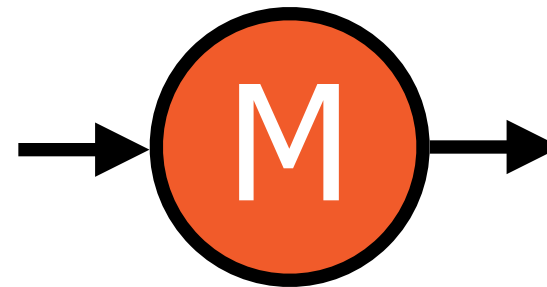
Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on mobiles
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	mike
4	Type	sale
4	Content	Clothes sale



Type	Count
offer	1
sale	2
order	1

Notification Type Count

Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on mobiles
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	mike
4	Type	sale
4	Content	Clothes sale



<offer, 1>

<sale, 1>

<order, 1>

<sale, 1>

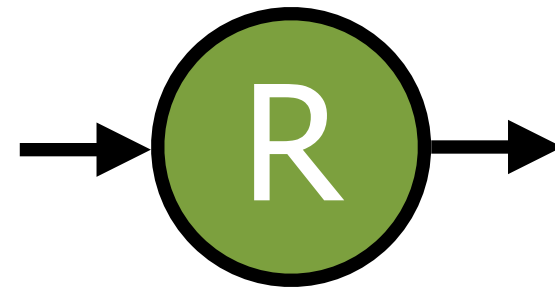
Notification Type Count

<offer, 1>

<sale, 1>

<order, 1>


<sale, 1>



Type	Count
offer	1
sale	2
order	1

Notification Type Count

Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on mobiles
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	mike
4	Type	sale
4	Content	Clothes sale



Type	Count
offer	1
sale	2
order	1

Implementing a MapReduce Job



Map

**A class where the
map logic is
implemented**

Reduce

**A class where the
reduce logic is
implemented**

Main

**A driver program
that sets up the job**

Implementing a MapReduce Job



The diagram consists of three colored squares arranged horizontally. The first square is purple and contains the word 'Map'. The second square is light green and contains the word 'Reduce'. The third square is light blue and contains the word 'Main'. Below each square is a descriptive text block.

Map

**A class where the
map logic is
implemented**

Reduce

A class where the
reduce logic is
implemented

Main

A driver program
that sets up the job

The Map Step

Map Class

Mapper Class

**The map logic is
implemented in a
class that extends the
Mapper Class**

The Map Step

Map Class

<input key type,
input value type,
output key type,
output value type>

Mapper Class

**This is a generic
class, with 4
type parameters**

Implementing a MapReduce Job



Map

A class where the
map logic is
implemented

Reduce

**A class where the
reduce logic is
implemented**

Main

A driver program
that sets up the job

The Reduce Step

Reduce Class

Reducer Class

**The reduce logic is
implemented in a
class that extends the
Reducer Class**

The Reduce Step

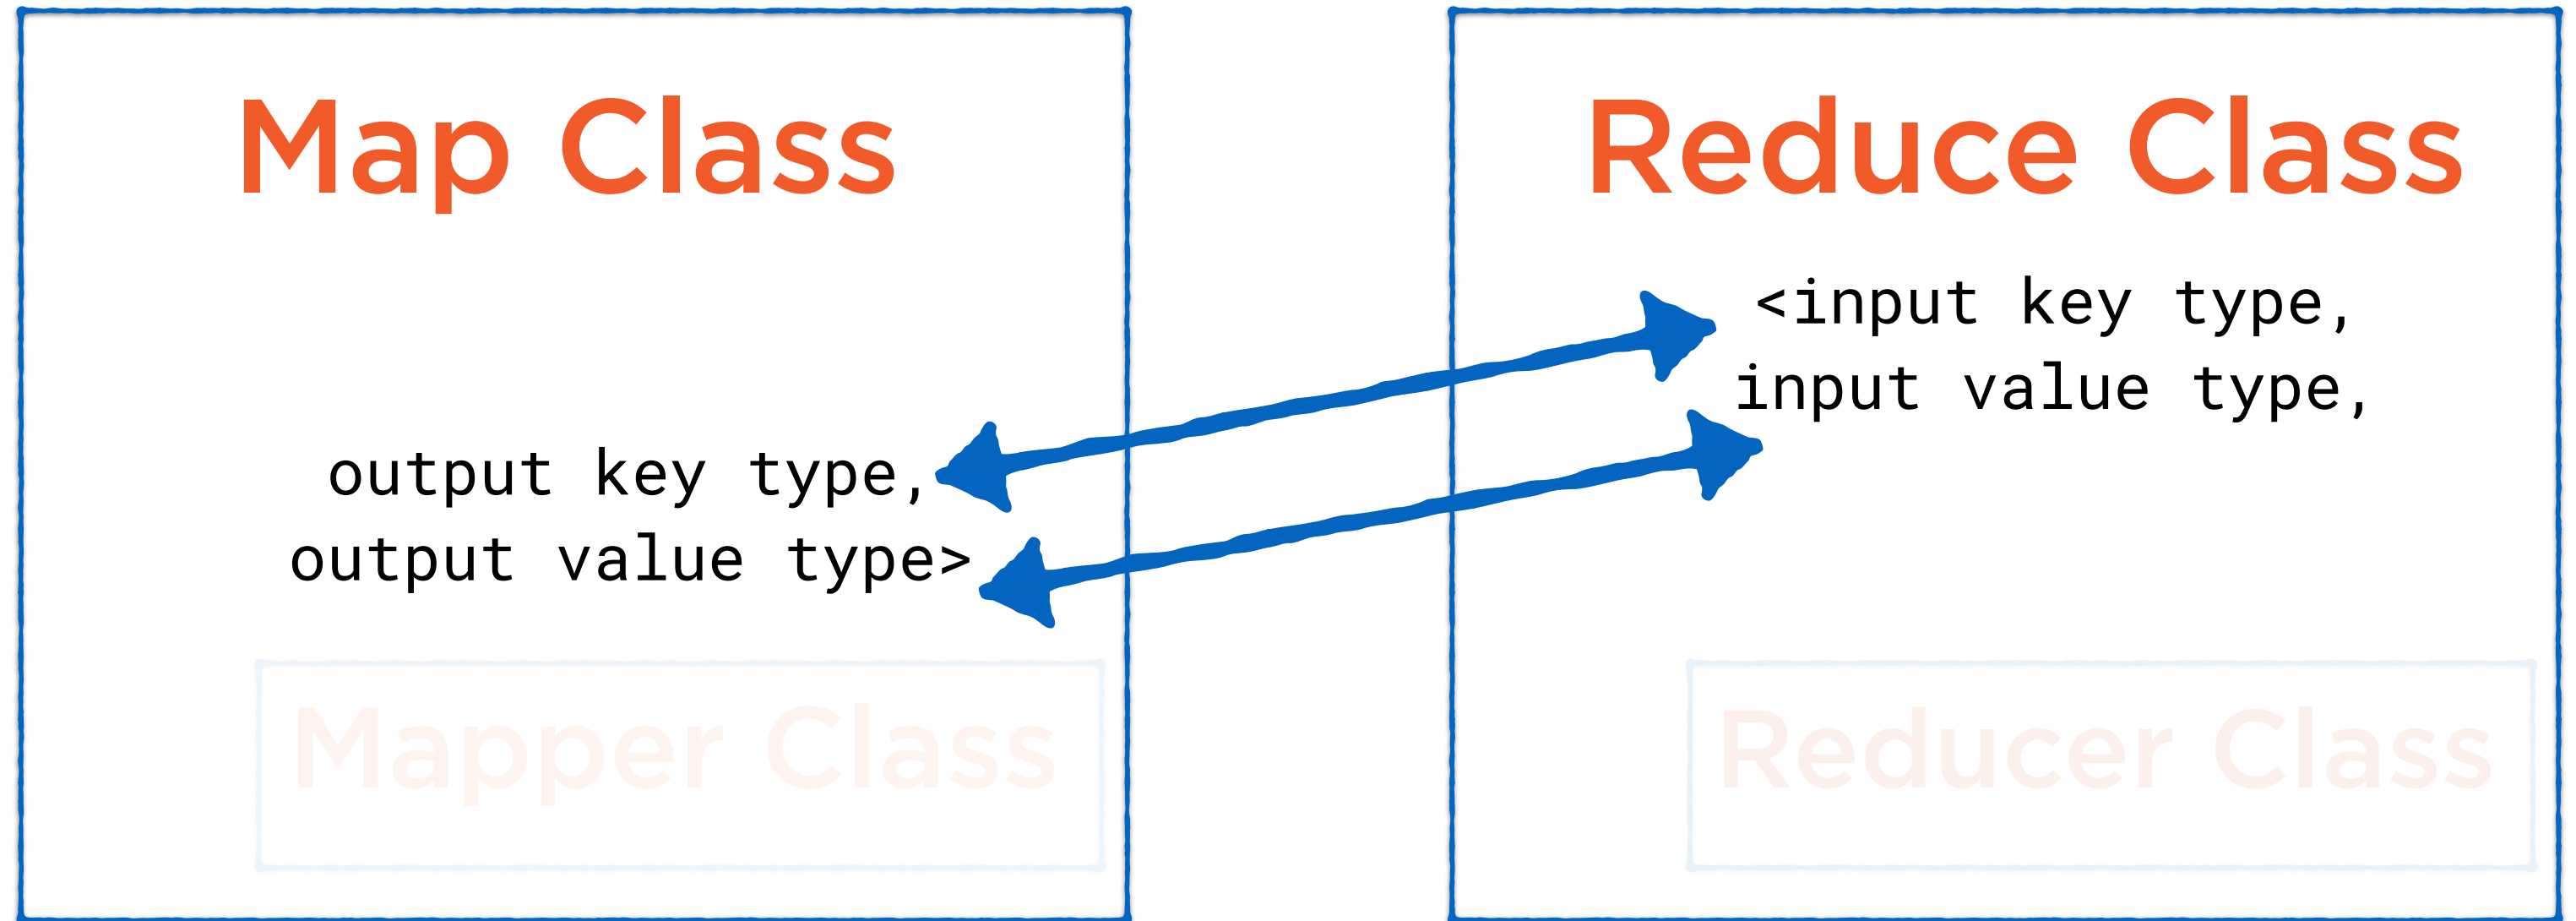
Reduce Class

<input key type,
input value type,
output key type,
output value type>

Reducer Class

**This is also a
generic class, with
4 type parameters**

Matching Data Types



Implementing a MapReduce Job



The diagram consists of three colored squares arranged horizontally. The first square is light purple and contains the word 'Map'. The second square is light green and contains the word 'Reduce'. The third square is blue and contains the word 'Main'. Below each square is a descriptive text block.

Map

A class where the
map logic is
implemented

Reduce

A class where the
reduce logic is
implemented

Main

**A driver program
that sets up the job**

Setting up the Job

The Mapper and Reducer classes are used by a Job that is configured in the Main Class



Main Class
Job Object

The diagram consists of two nested rectangles with blue borders. The outer rectangle is labeled 'Main Class' in orange text at its top right corner. Inside it, a smaller rectangle is labeled 'Job Object' in orange text at its top right corner. This visualizes that the Job Object is contained within the Main Class.

Setting up the Job

The Job has a bunch of properties that need to be configured

Main Class

Job Object

Input filepath

Output filepath

Mapper class

Reducer class

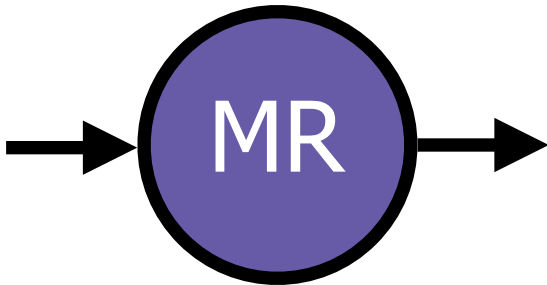
Output data types

Demo

Set up a MapReduce job which runs on the census HBase table

Marital Status Count

Id	Column	Value
1	name	mike
1	marital_status	married
1	employed	yes
2	name	john
2	marital_status	divorced
2	field	real estate
3	name	jill
3	marital_status	married
3	employed	yes
4	name	ben
4	marital_status	separated
4	education_level	undergrad



Marital Status	Count
married	2
divorced	1
separated	1

Marital Status Count


Id	Column	Value
1	name	mike
1	marital_status	married
1	employed	yes
2	name	john
2	marital_status	divorced
2	field	real estate
3	name	jill
3	marital_status	married
3	employed	yes
4	name	ben
4	marital_status	separated
4	education_level	undergrad



<married, 1>
<divorced, 1>
<married, 1>
<separated, 1>

Marital Status Count

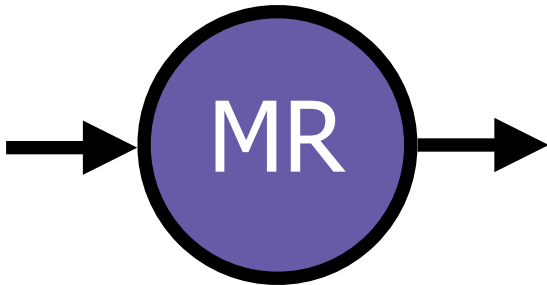
<married, 1>
<divorced, 1>
<married, 1>
<separated, 1>



Marital Status	Count
married	2
divorced	1
separated	1

Marital Status Count

Id	Column	Value
1	name	mike
1	marital_status	married
1	employed	yes
2	name	john
2	marital_status	divorced
2	field	real estate
3	name	jill
3	marital_status	married
3	employed	yes
4	name	ben
4	marital_status	separated
4	education_level	undergrad



Marital Status	Count
married	2
divorced	1
separated	1

Summary

**A brief overview of the Hadoop
MapReduce framework**

**Ran MapReduce using HBase tables as
a source and a sink of data**