Why do we need NoSQL / HBase?

Relational Model

- Pros
  - **Standard** persistence model
  - **Transactions handle**
    - **concurrency , consistency**
  - efficient and robust structure for storing data

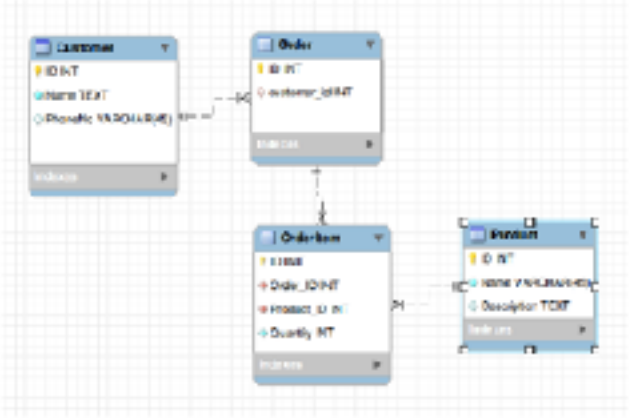

Relational Databses vs HBase - Scaling

RDBMS - Scale UP approach

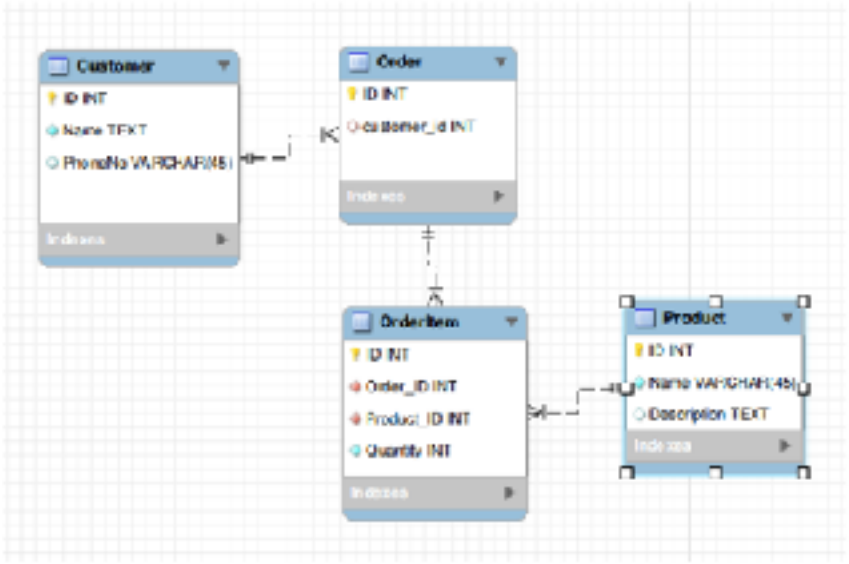

Vertical scale = **big box**

RDBMS Scaling Up Example - eBay

- Back End Oracle Database server **scaled vertically to a larger machine** (Sun E10000)

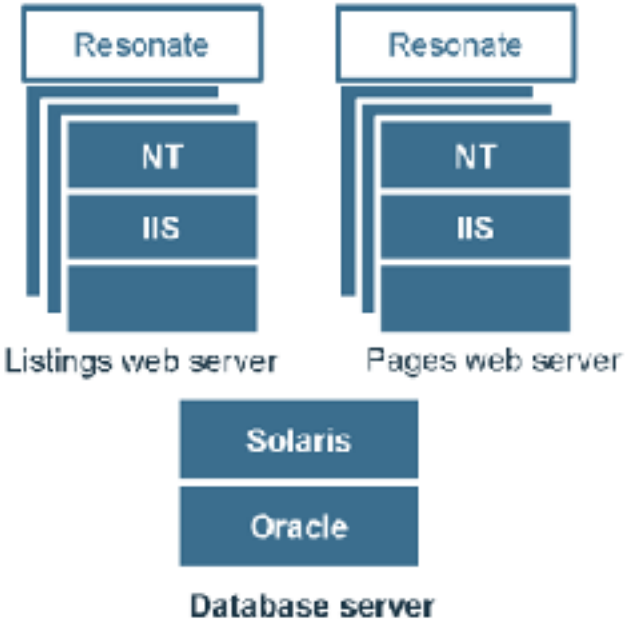

http://www.addsimplicity.com/downloads/eBaySDForum2006-11-29.pdf

What changed to bring on NoSQL?

Lots of data, the need to scale horizontally

Horizonal scale: split table by rows into partitions across a cluster

Key	colB	colC
val	val	val
xxx	val	val

id 1-1000

Key	colB	colC
val	val	val
xxx	val	val

id 1000-2000

Key	colB	colC
val	val	val
xxx	val	val

id 2000=3000

- **Horizontal** scaling
  - Cheaper than vertical
  - parallel execution
- Relational databases were **not designed to do this** automatically

Facebook 2010

- 9000 memcache instances
- **4000 Shards mysql**

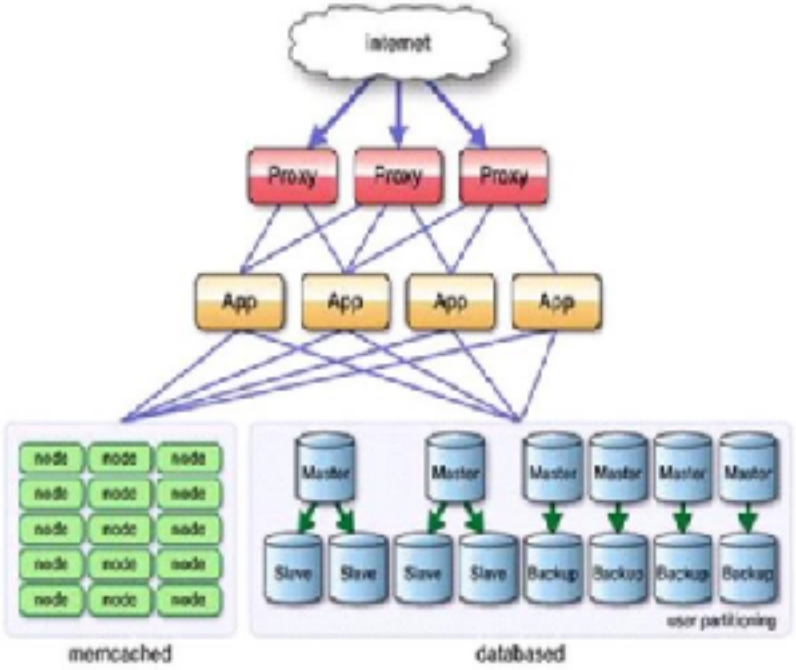

http://gigaom.com/2011/07/07/facebook-trapped-in-mysql-fate-worse-than-death/

Relational Databases vs. HBase – Data Storage Model

RDBMS

Storage Model

HBase

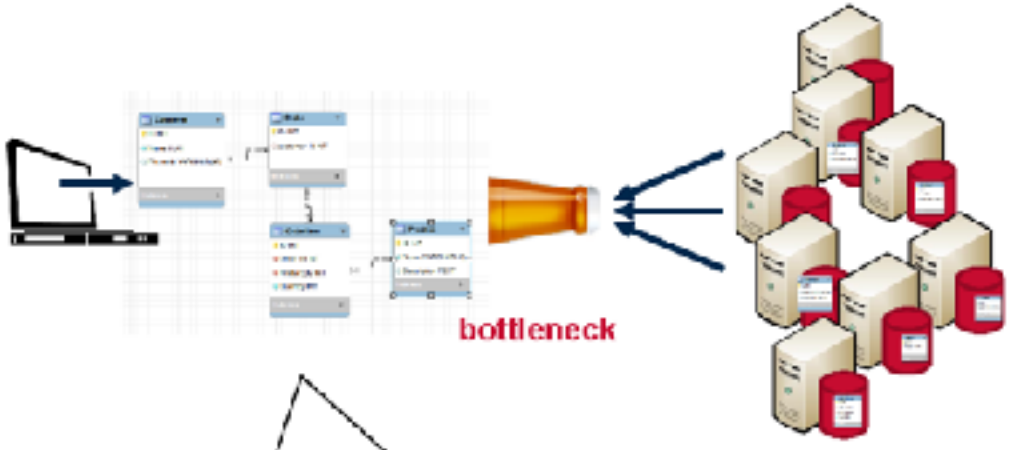

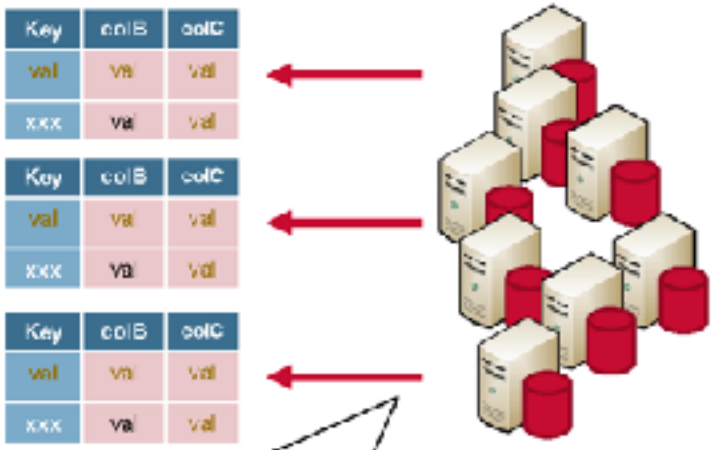

Distributed Joins, Transactions **do not scale**

Data that is accessed together is stored together

Hbase designed for Distribution, Scale, Speed

Big Table

- Distributed Storage System
- Paper published in 2006.

Google File System

MapReduce

Runs on commodity hardware
Designed to Scale

HBase is a ColumnFamily oriented Database

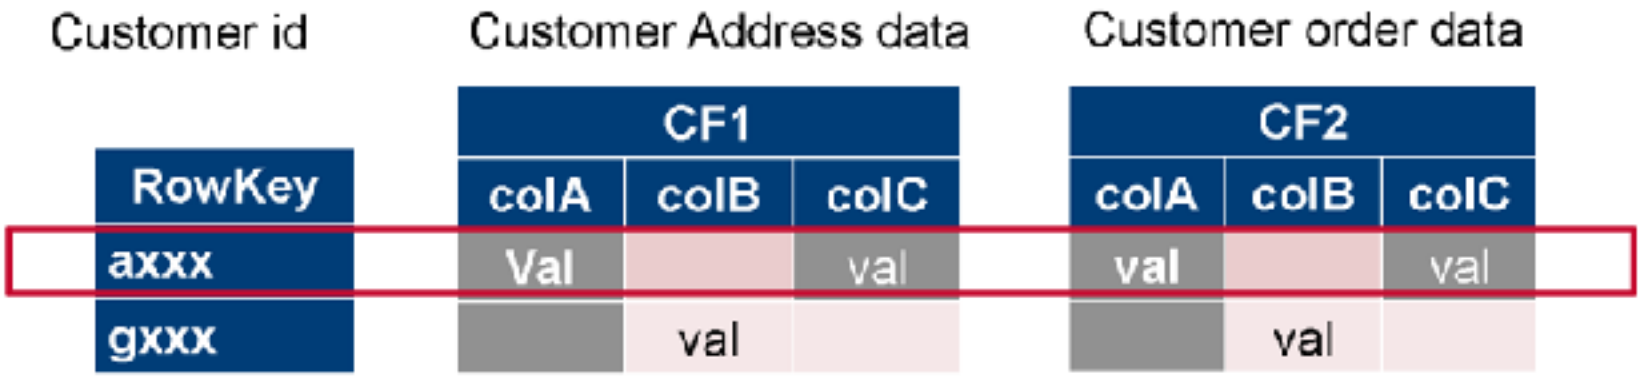

Data is accessed and stored together:

- RowKey is the primary index
- Column Families group similar data by row key

HBase is a Distributed Database

Put, Get by Key

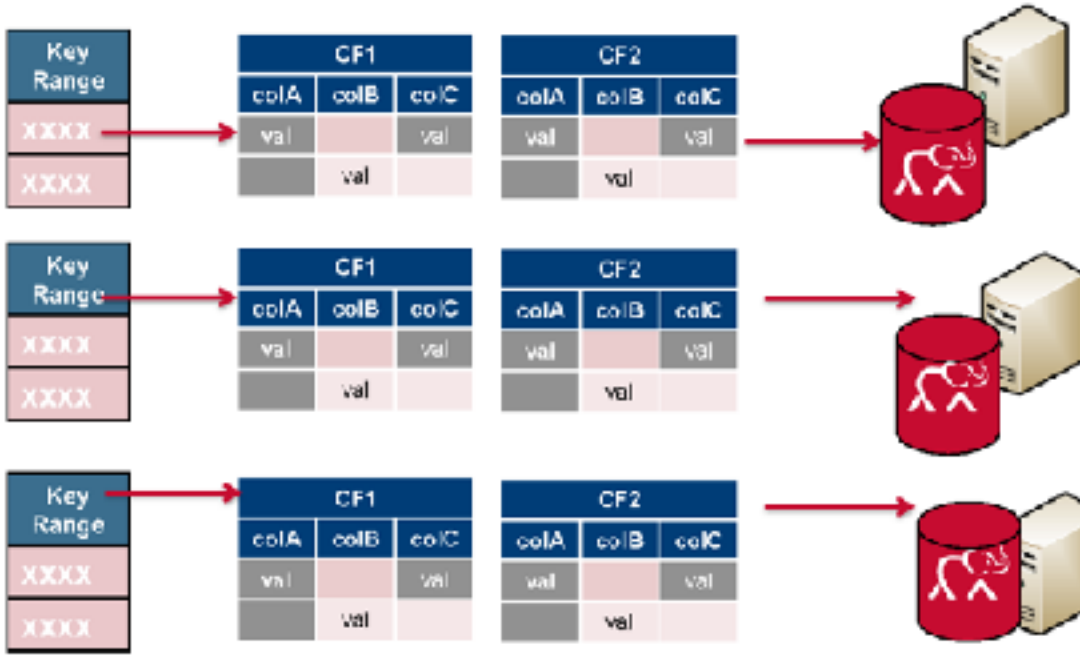

Data is automatically distributed across the cluster

- **Key range** is used for **horizontal partitioning**

Column Family Databases

distributed data stored and accessed together:

- Pros
  - **scalable**
  - **Fast Writes and Reads**
- Cons
  - **No** joins
  - No dynamic queries
  - Need to **know** how data will be **queried in advance**

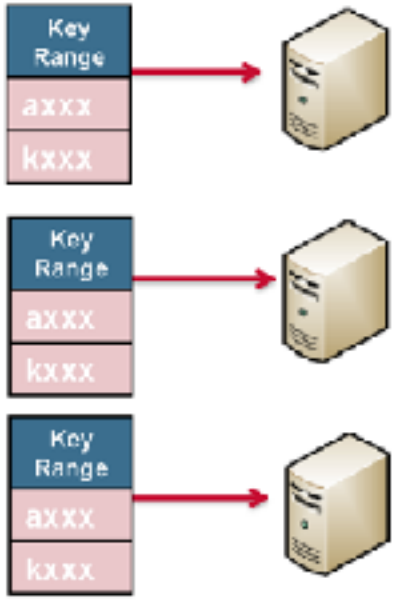

Hbase Data Model

HBase Data Model- Row Keys

RowKey	Address			Order				...
	street	city	state	Date	ItemNo	Ship Address	Cost	
smithj	val		val	val			val	
spata								
sxxxx	val			val	val	val		
...								
turnerb	val	val	val	val	val	val	val	
...								
	val							
twistr	val		val	val			val	
...								
zaxx	val	val	val	val	val	val		
zxxx	val						val	

Row Keys: identify the rows in an HBase table.

Tables are split into Regions = contiguous keys

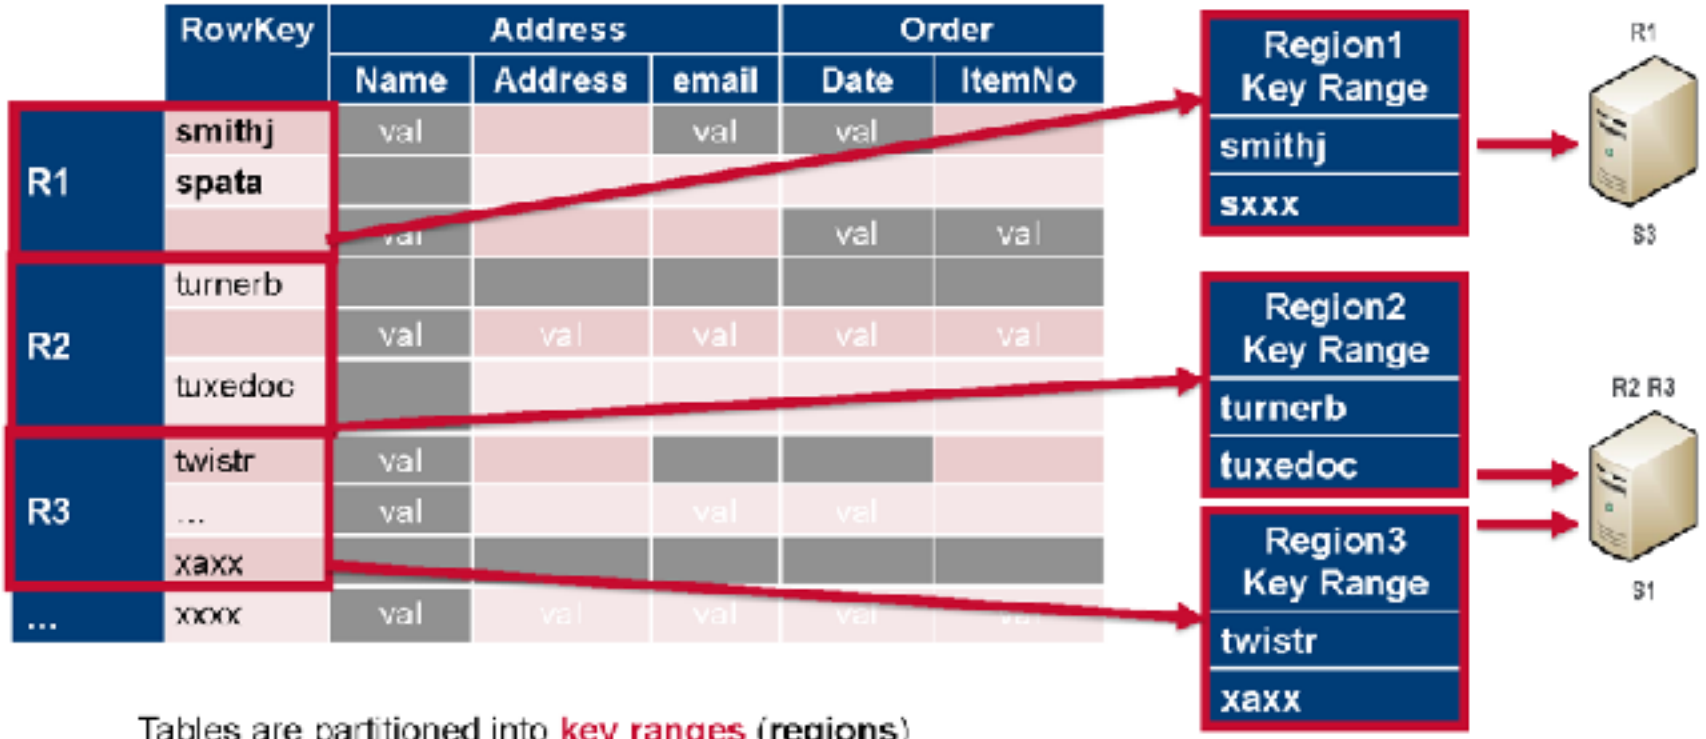

Tables are partitioned into **key ranges** (**regions**)
Region= served by nodes (Region Servers)
Regions are spread across cluster

Column Families are stored Separately

- column families are stored and can be accessed separately

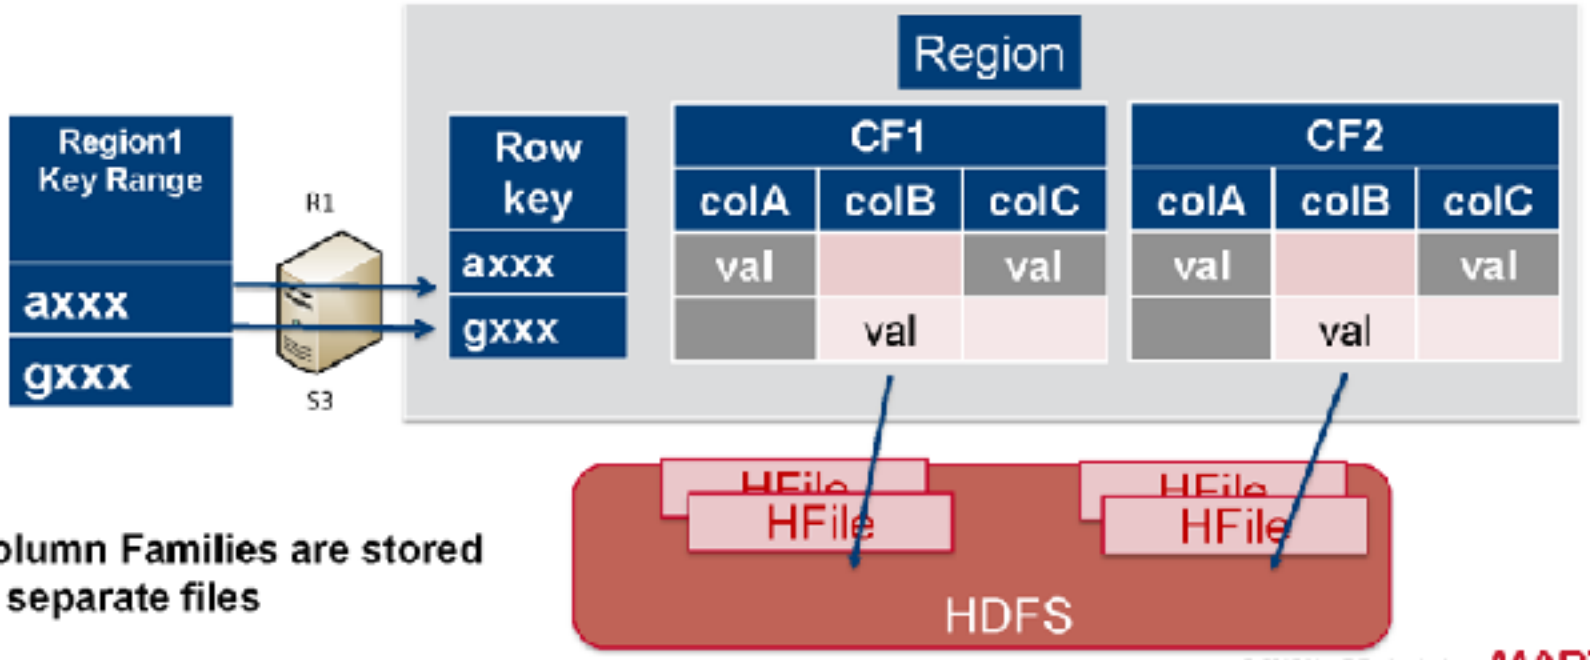

Column Families are stored in separate files

HBase Data Model - Cells

- Data is stored in **Key Value** format
- Value for each **cell** is specified by complete coordinates:
  - (Row key, ColumnFamily, Column Qualifier, timestamp) => Value

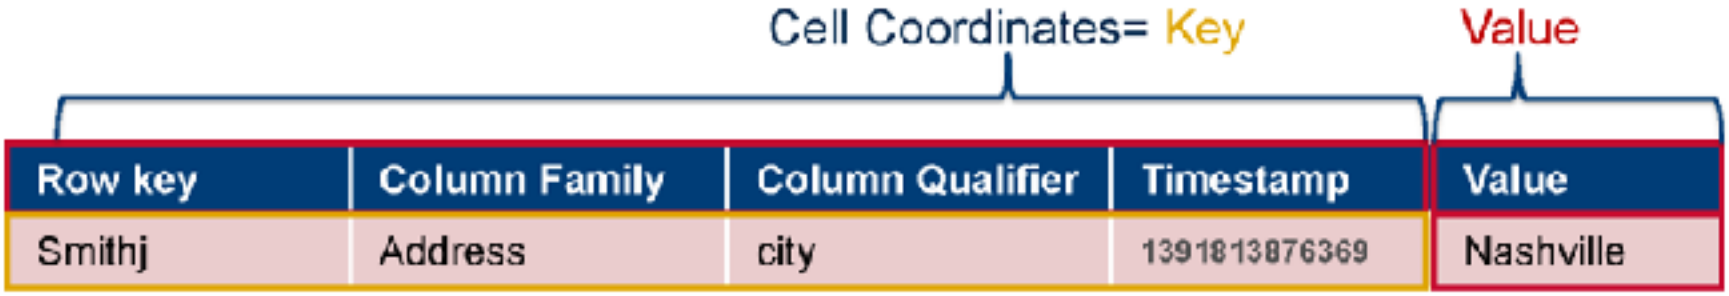

Logical Data Model vs Physical Data Storage

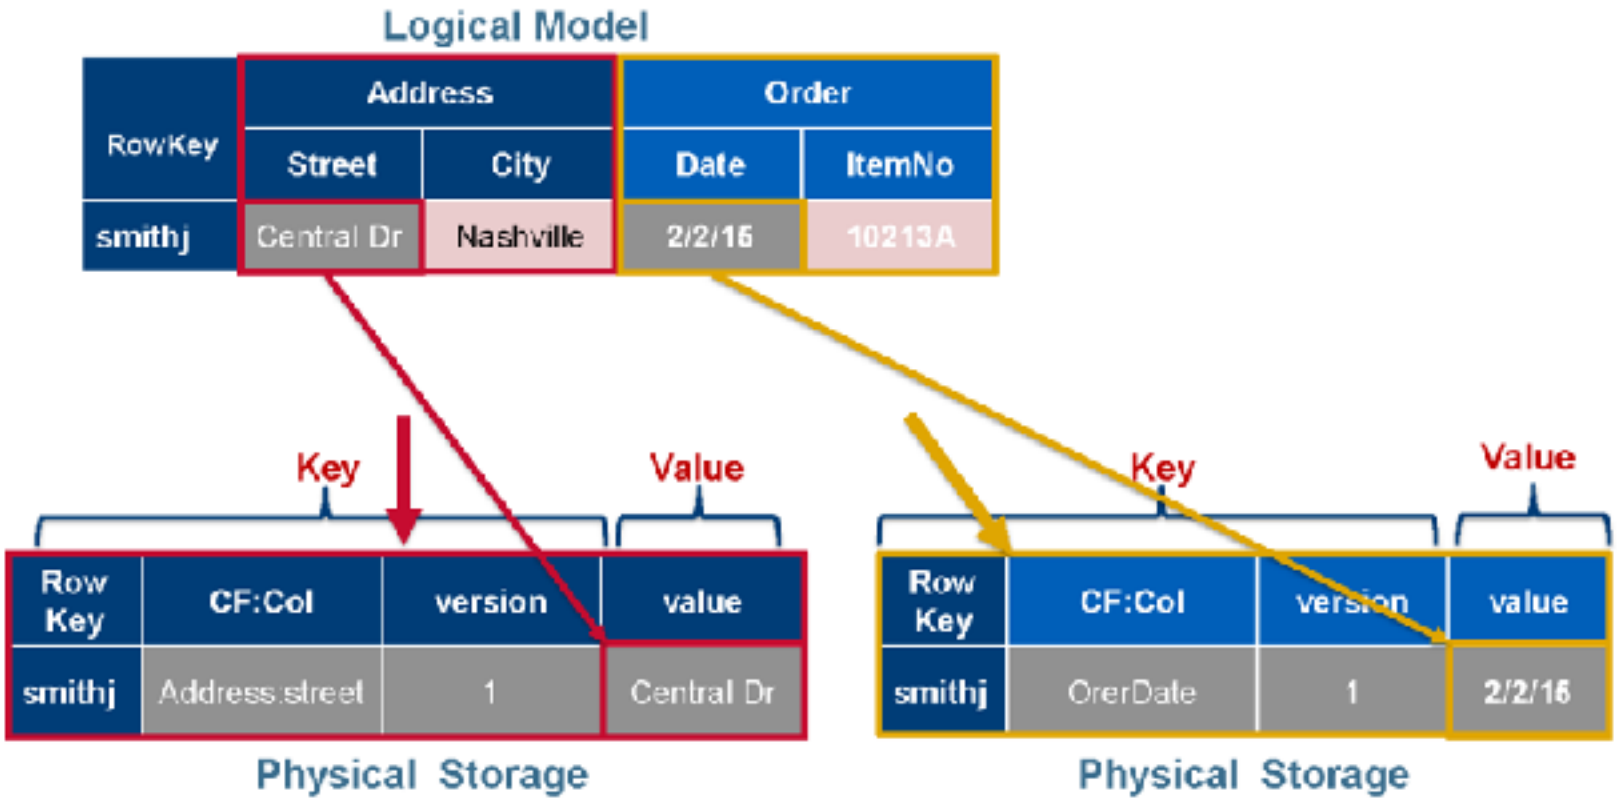

Sparse Data with Cell Versions

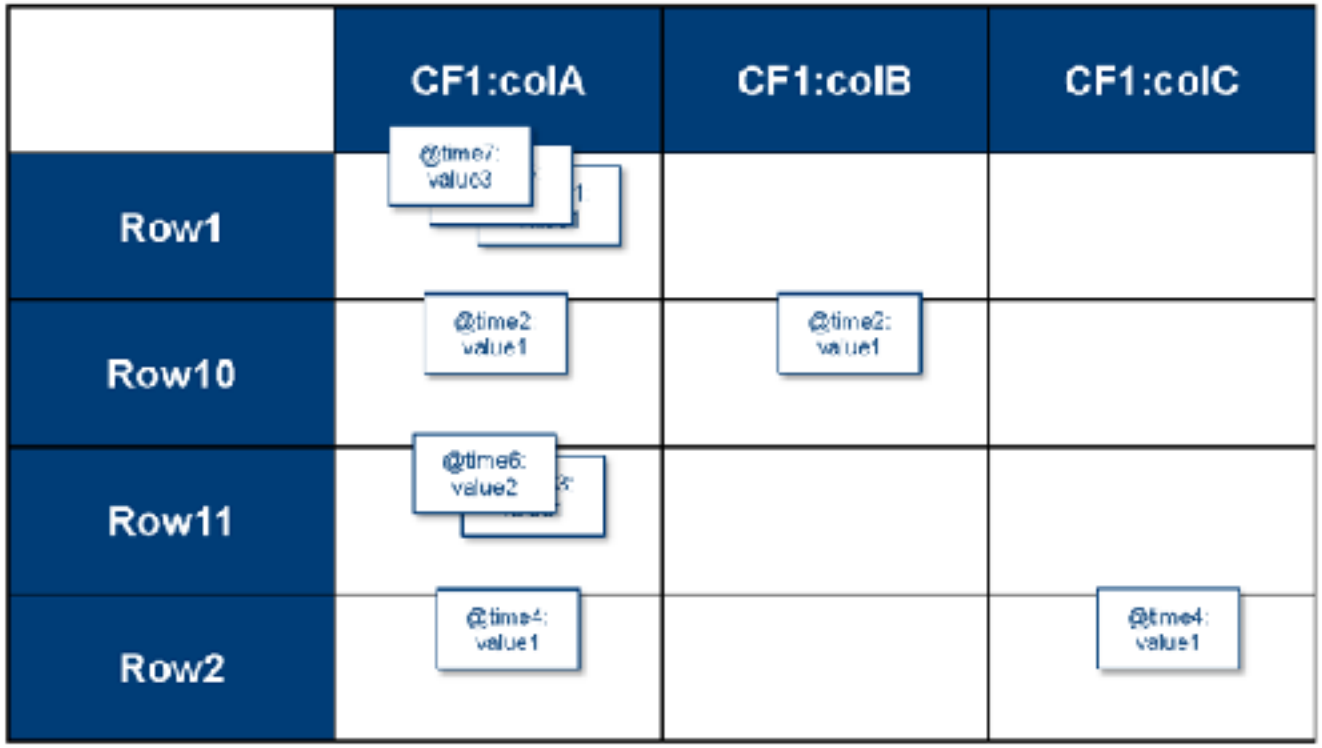

Versioned Data

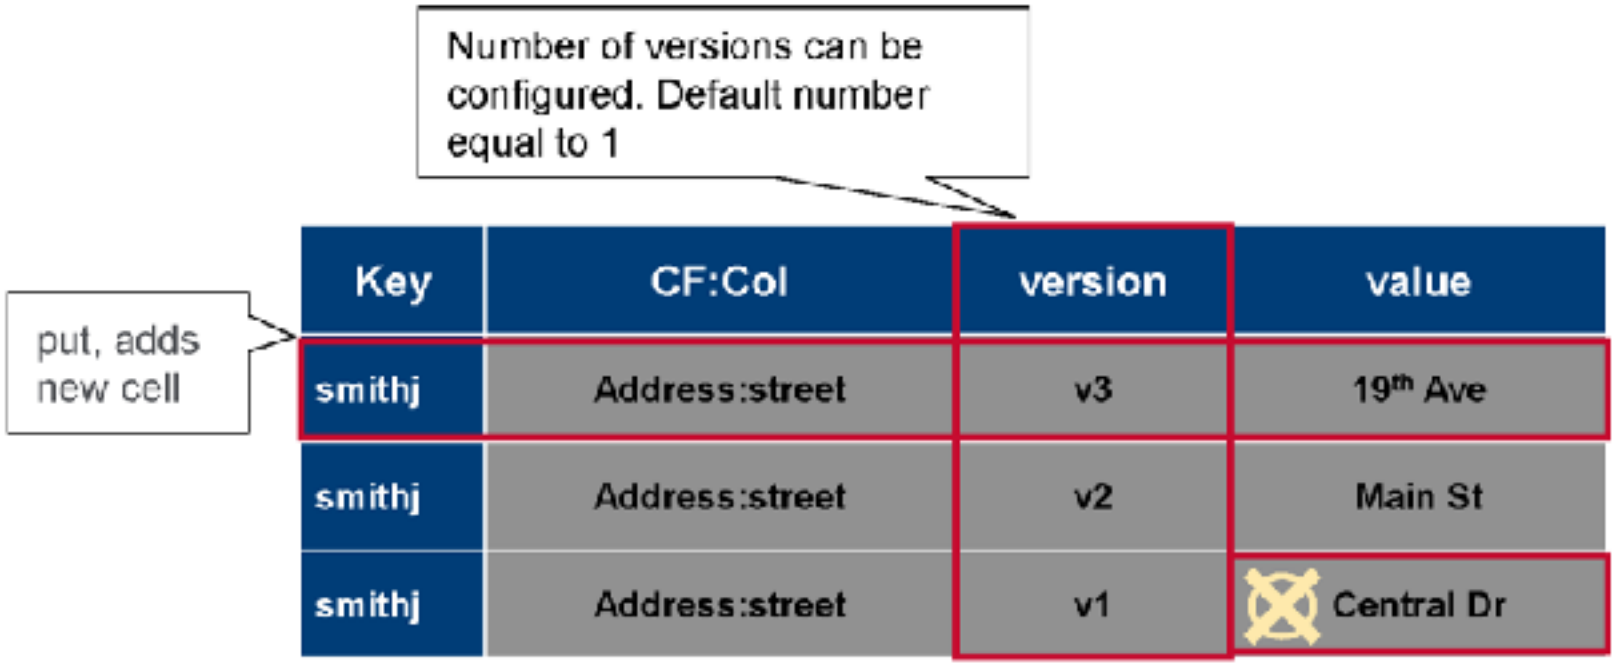

Logical Data Model vs Physical Data Storage

Row Key	CF1		CF2	
	ca	cb	ca	cd
ra	1		2	
rb		3,4		
rc	5		6,7	8

Logical Model

Column families are stored separately
Row keys, Qualifiers are sorted **lexicographically**

Physical Storage

Key | Value

Key	CF1:Col	version	value
ra	cf1:ca	1	1
rb	cf1:cb	2	4
rb	cf1:cb	1	3
rc	cf1:ca	1	5

Key | Value

Key	CF2:Col	version	value
ra	cf2:ca	1	2
rc	cf2:ca	2	7
rc	cf2:ca	1	6
rc	cf2:cd	1	8

HBase Table is a Sorted Map of Maps

In Java this is the table:

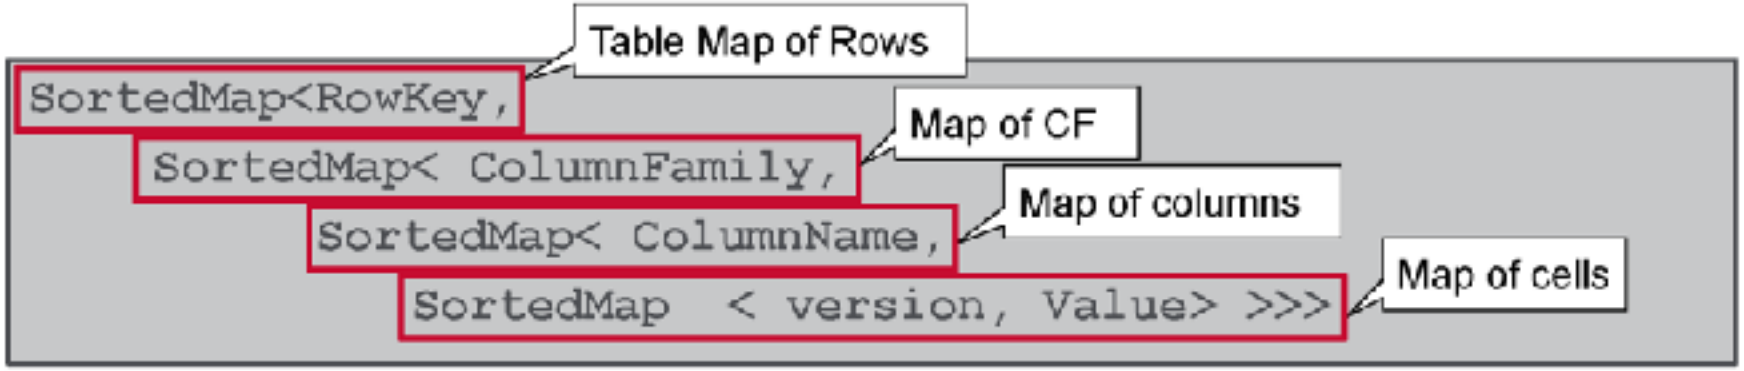

Key	CF:Col	version	value
smithj	Address:street	v2	Main St
smithj	Address:street	v1	Central Dr
spata	Address:street	v1	High Ave
turnerb	Address:street	v1	Cedar St

Key	CF:Col	version	value
smithj	Order:Date	v1	2/2/15
spata	Order:Date	v1	1/31/14
turnerb	OrderDate	v1	7/8/14

HBase Table - SortedMap

```
<smithj,<Address, <street, <v1, Central Dr>>
                  <street, <v2, Main St>>
              <Order <Date, <v1, 2/2/15>>>
<spata,<Address, <street, <v1,High Ave>>
              <Order <Date, <v1, 1/31/14>>>
<turnerb,<Address, <street, <v1,Cedar St>>
              <Order <Date, <v1, 7/8/14>>>
```

Key	CF:Col	version	value
smithj	Address:street	v2	Main St
smithj	Address:street	v1	Central Dr
spata	Address:street	v1	High Ave
turnerb	Address:street	v1	Cedar St

Key	CF:Col	version	value
smithj	Order:Date	v1	2/2/15
shawa	Order:Date	v1	1/31/14
turnerb	OrderDate	v1	7/8/14

Basic Table Operations

- **Create Table, define Column Families before data is imported**
  - but not the rows keys or number/names of columns
- Low level API, technically more demanding
- **Basic data access operations (CRUD):**

put	Inserts data into rows (both create and update)
get	Accesses data from one row
scan	Accesses data from a range of rows
delete	Delete a row or a range of rows or columns

Create HBase Table – Using HBase Shell

```
hbase> create '/user/user01/Customer', {NAME =>'Address'} , {NAME =>'Order'}
hbase> put '/user/user01/Customer', 'smithj', 'Address:street', 'Central Dr'
hbase> put '/user/user01/Customer', 'smithj', 'Order:Date', '2/2/15'
hbase> put '/user/user01/Customer', 'spata', 'Address:city', 'Columbus'
hbase> put '/user/user01/Customer', 'spata', 'Order:Date', '1/31/14'
```

Row Key	Address			Order	
	street	city	state	Date	ItemNo
smithj	Central Dr	Nashville	TN	2/2/15	10213A
spata	High Ave	Columbus	OH	1/31/14	23401V
turnerb	Cedar St	Seattle	WA	7/8/14	10938A

Create HBase Table – Using HBase Shell

```
hbase> get '/user/user01/Customer', 'smithj'

hbase> scan '/user/user01/Customer'

hbase> describe '/user/user01/Customer'
```

HBase Architecture

Data flow for Writes, Reads

Designed to Scale

HBase Architectural Components

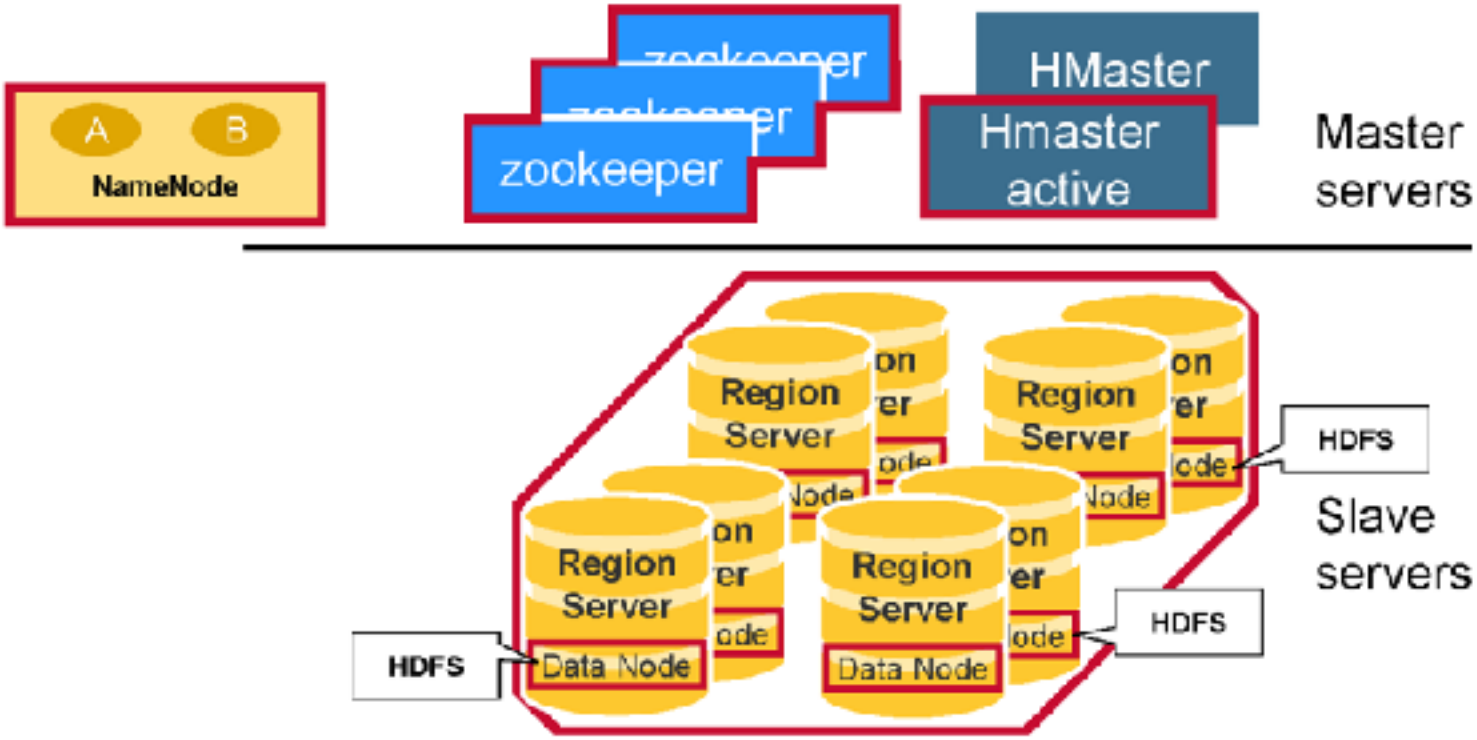

Regions

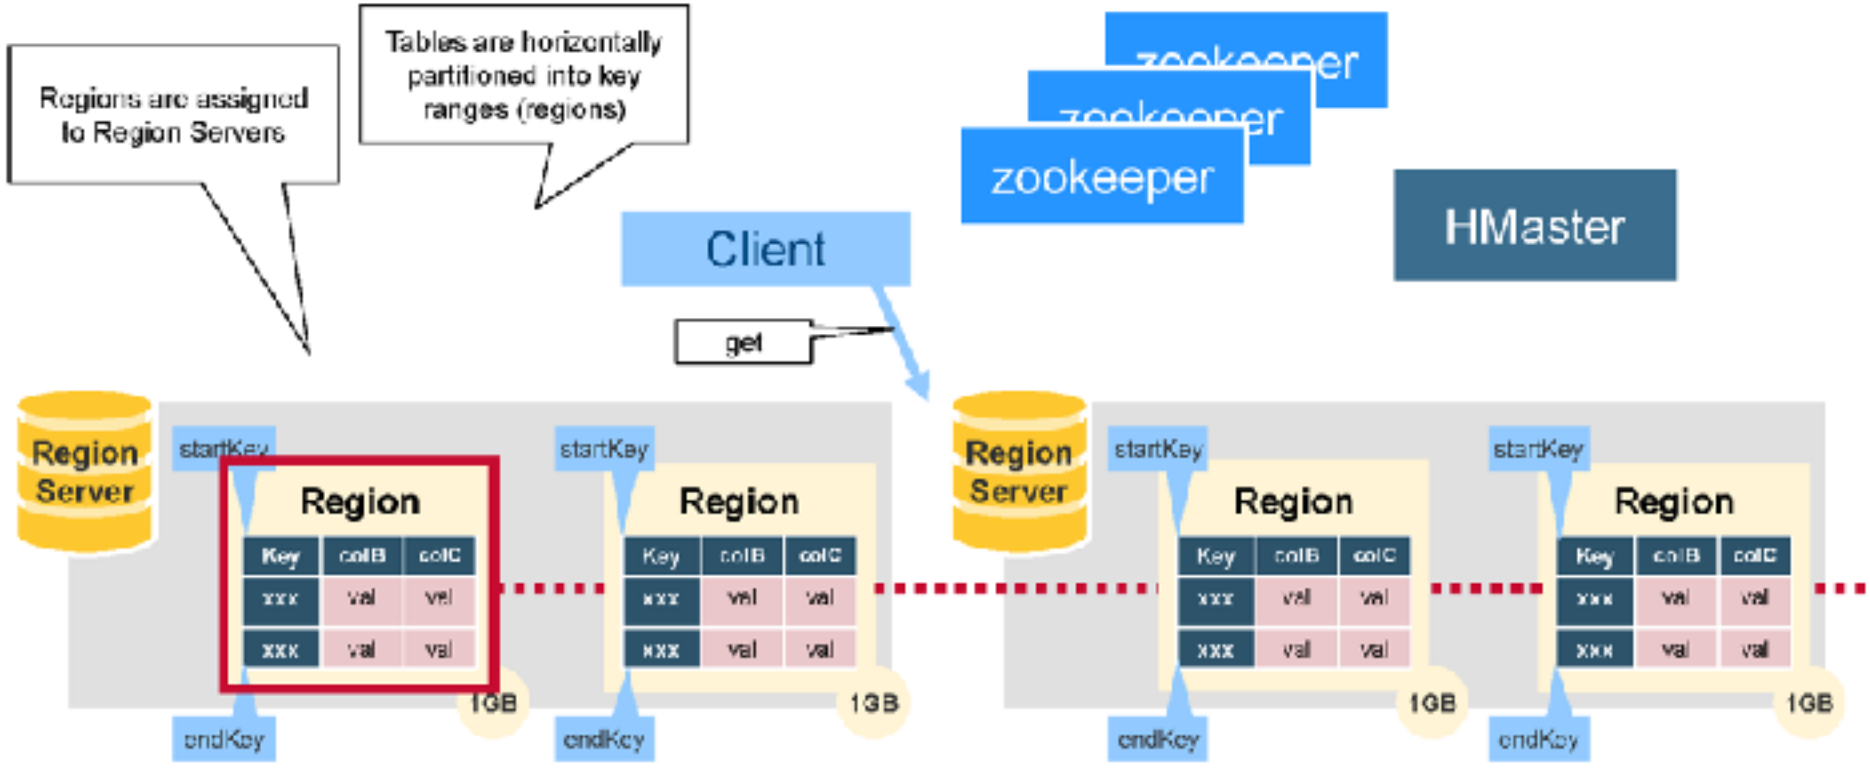

HBase HMaster

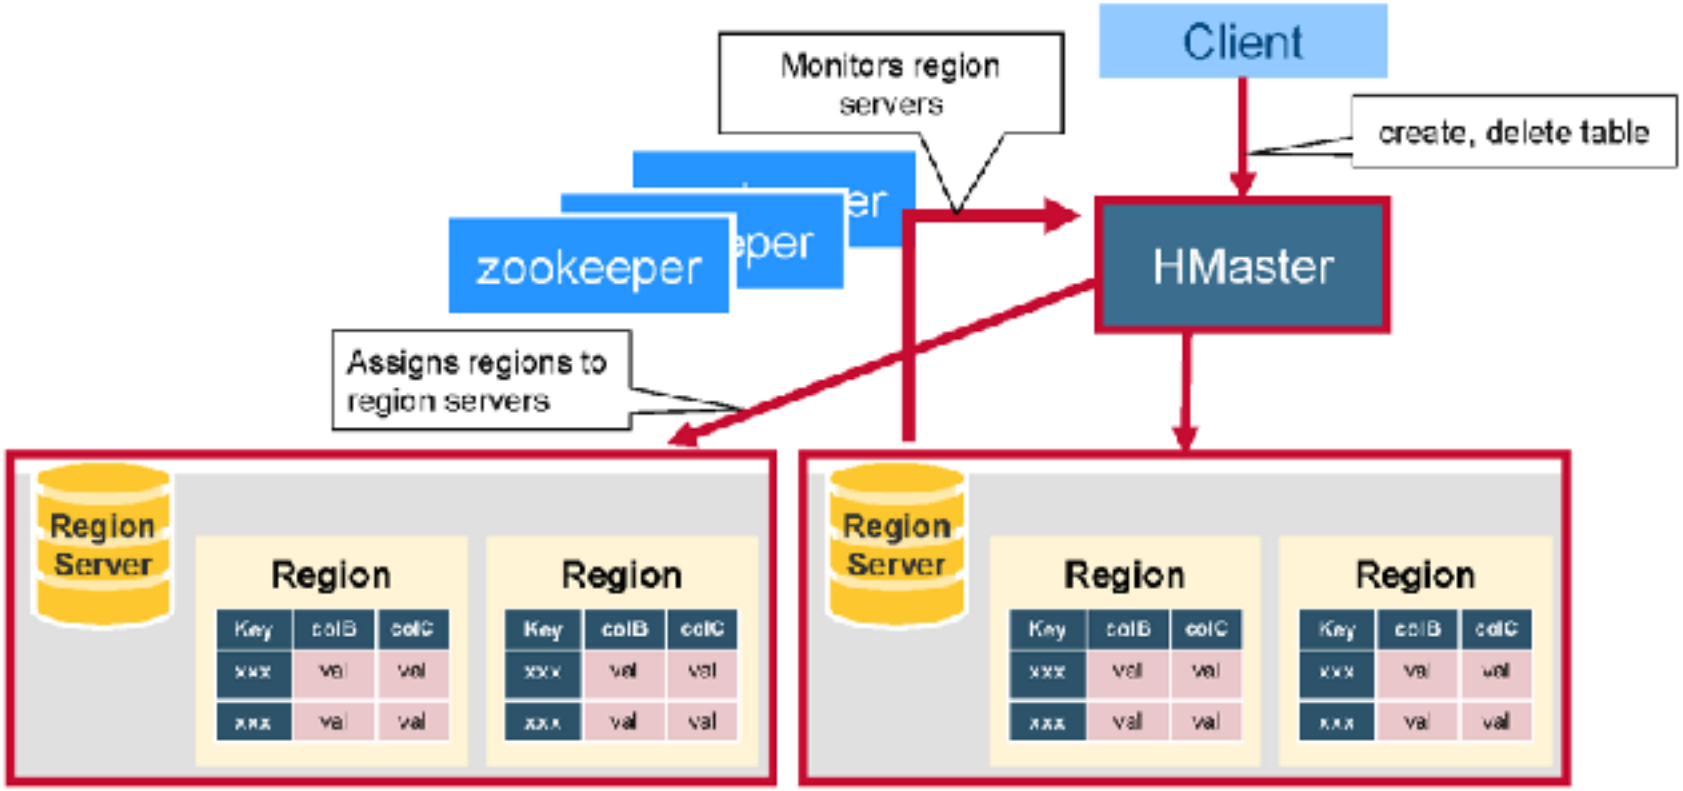

Zookeeper The Coordinator

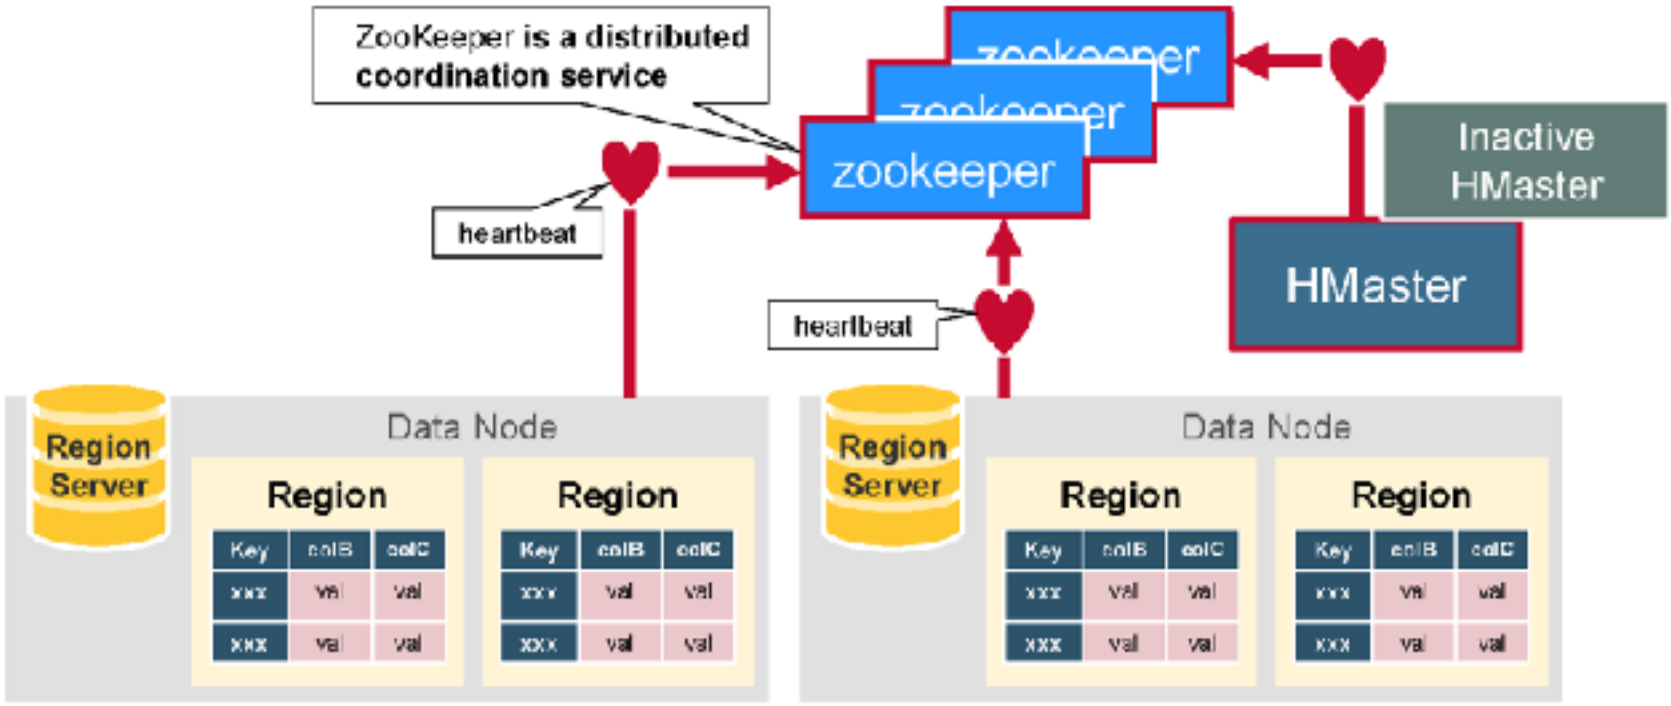

How the Components Work Together

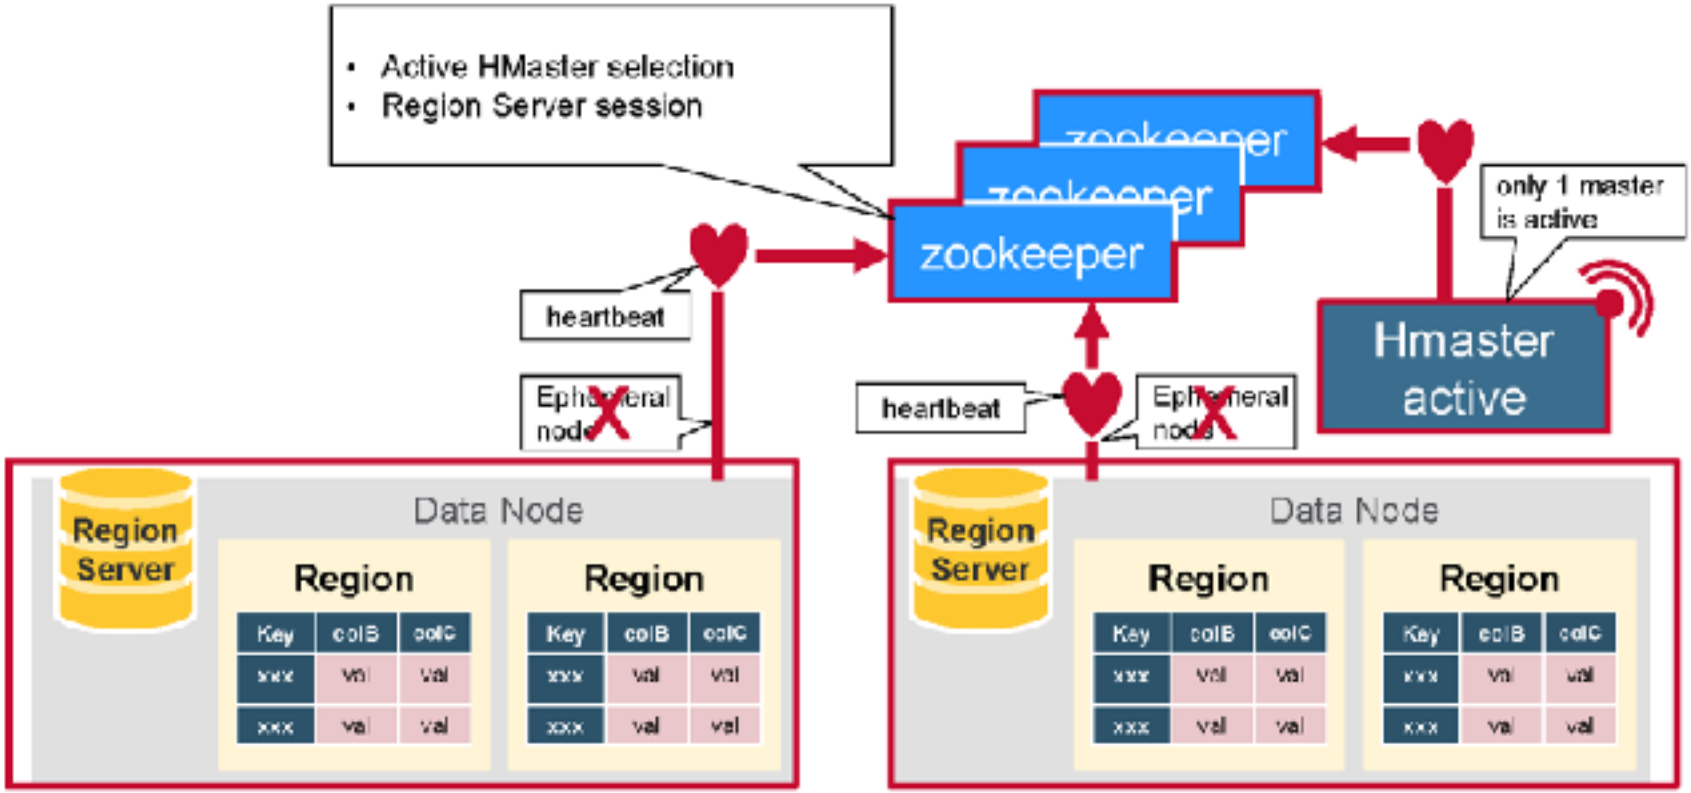

HBase First Read or Write

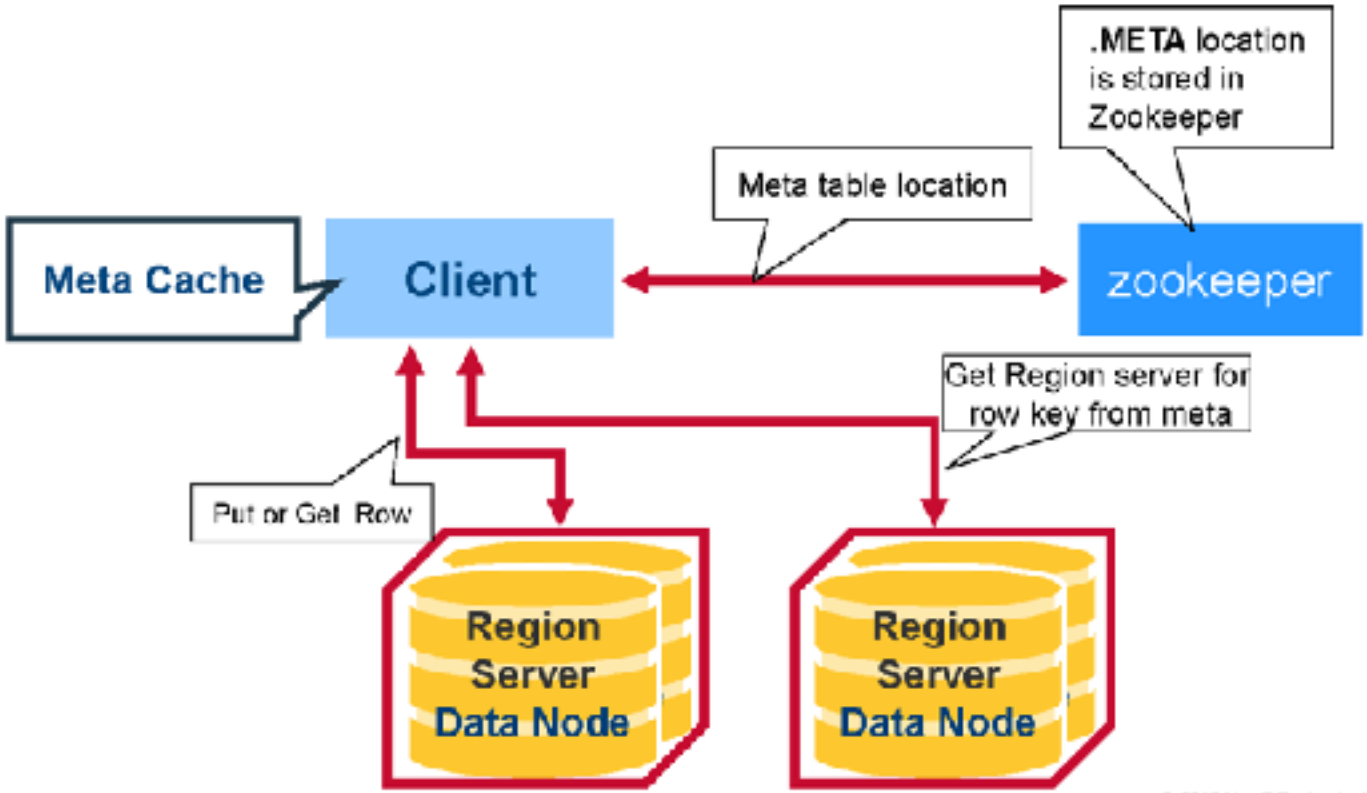

HBase Meta Table

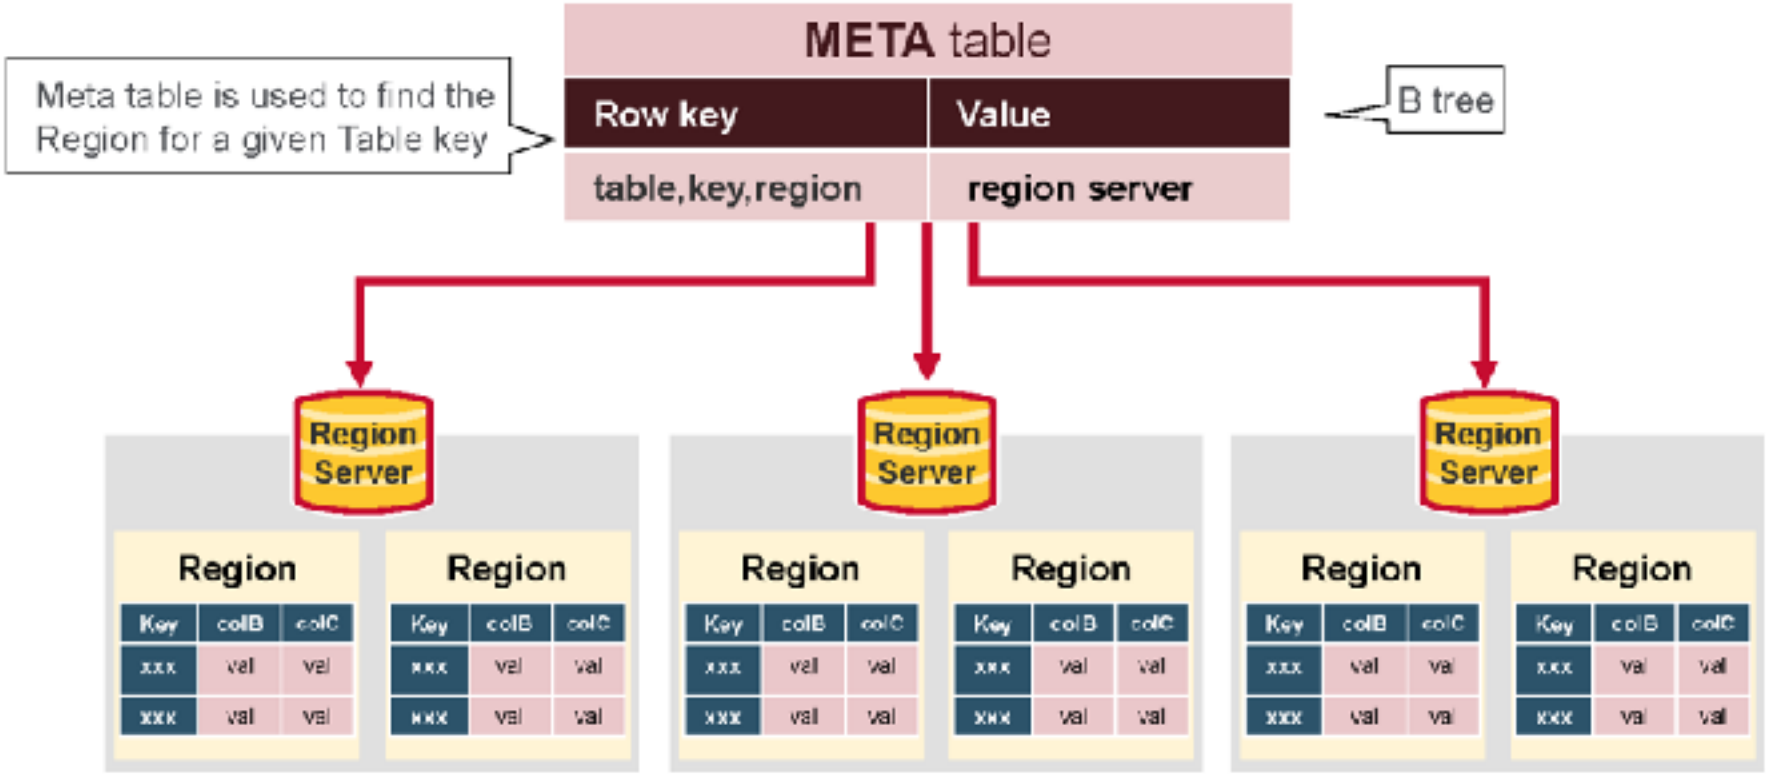

Region Server Components

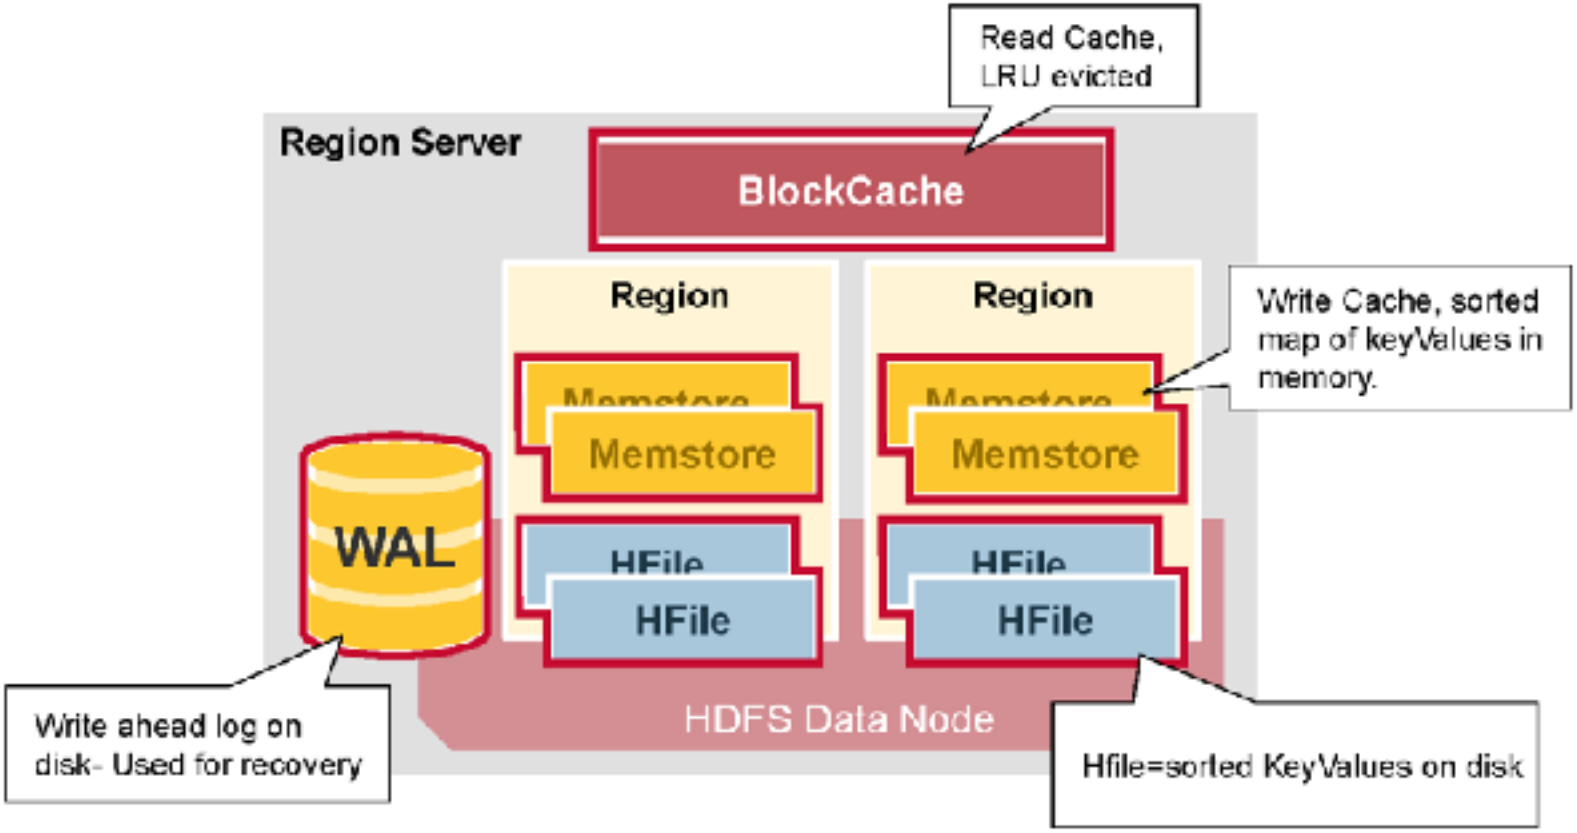

HBase Write Steps

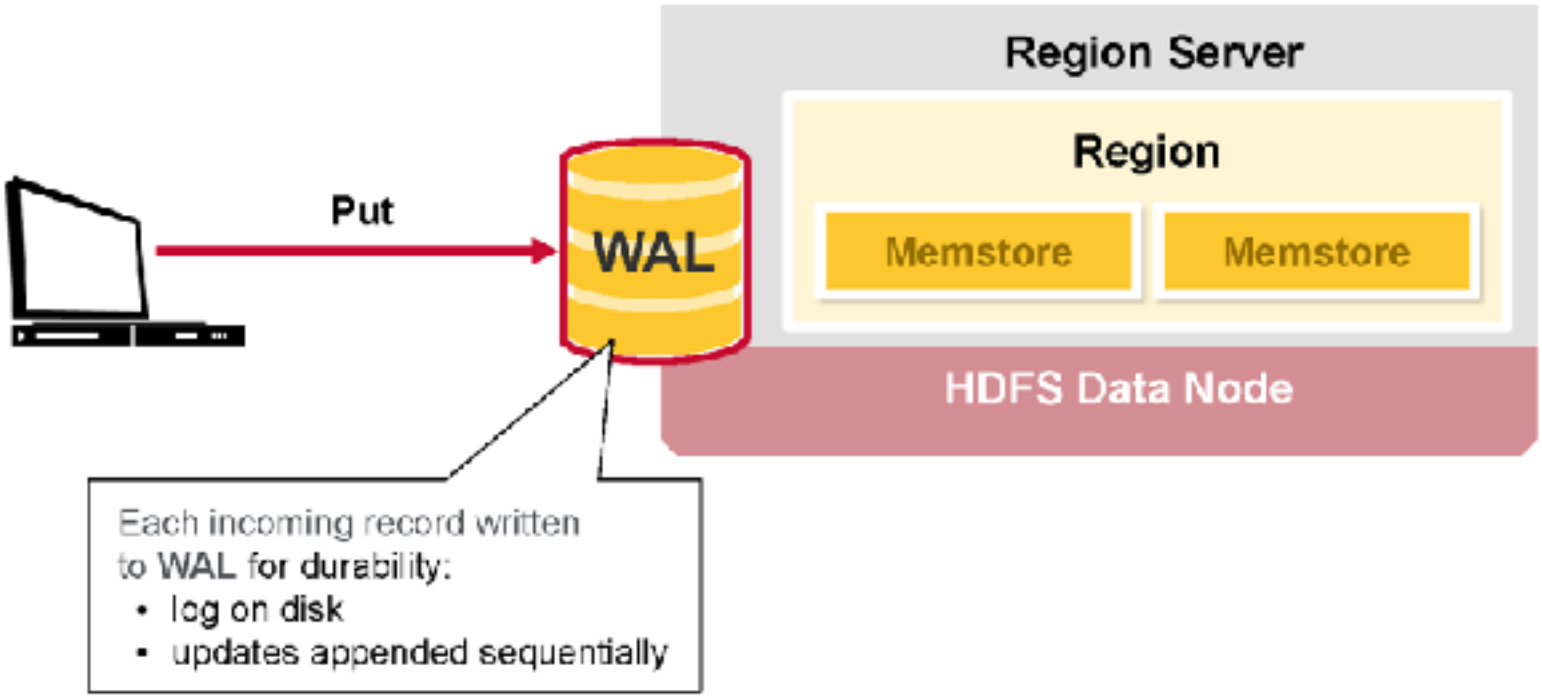

HBase Write Steps – (2)

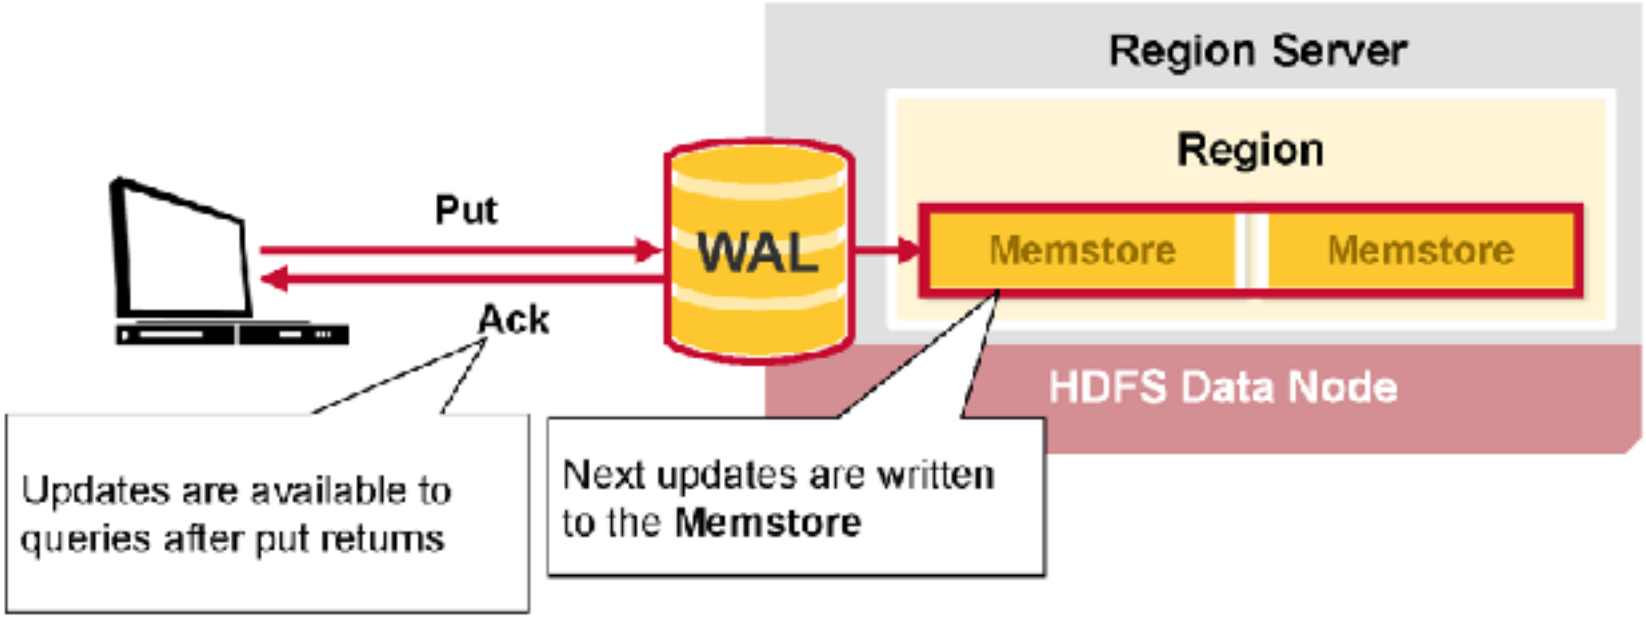

HBase Memstore

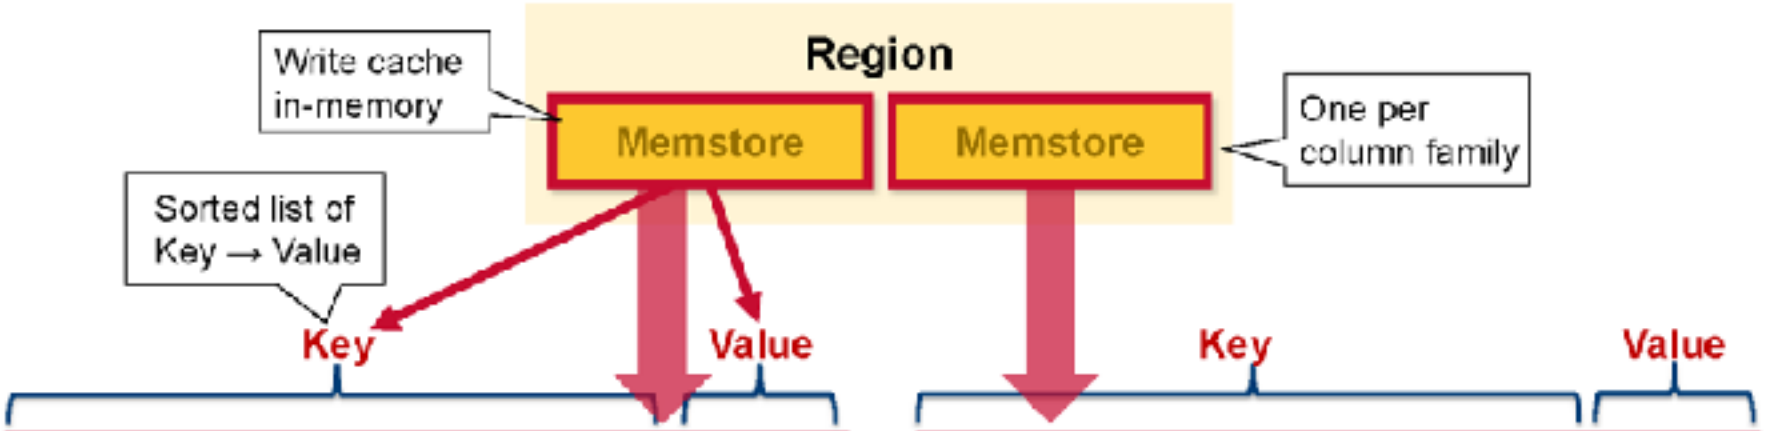

Key	CF1:Col	version	value
ra	cf1:ca	v1	1
rb	cf1:cb	v2	4
rb	cf1:cb	v1	3
rc	cf1:ca	v1	5

Key	CF2:Col	version	value
ra	cf2:ca	v1	2
rc	cf2:ca	v2	7
rc	cf2:ca	v1	6
rc	cf2:cd	v1	8

HBase Region Flush

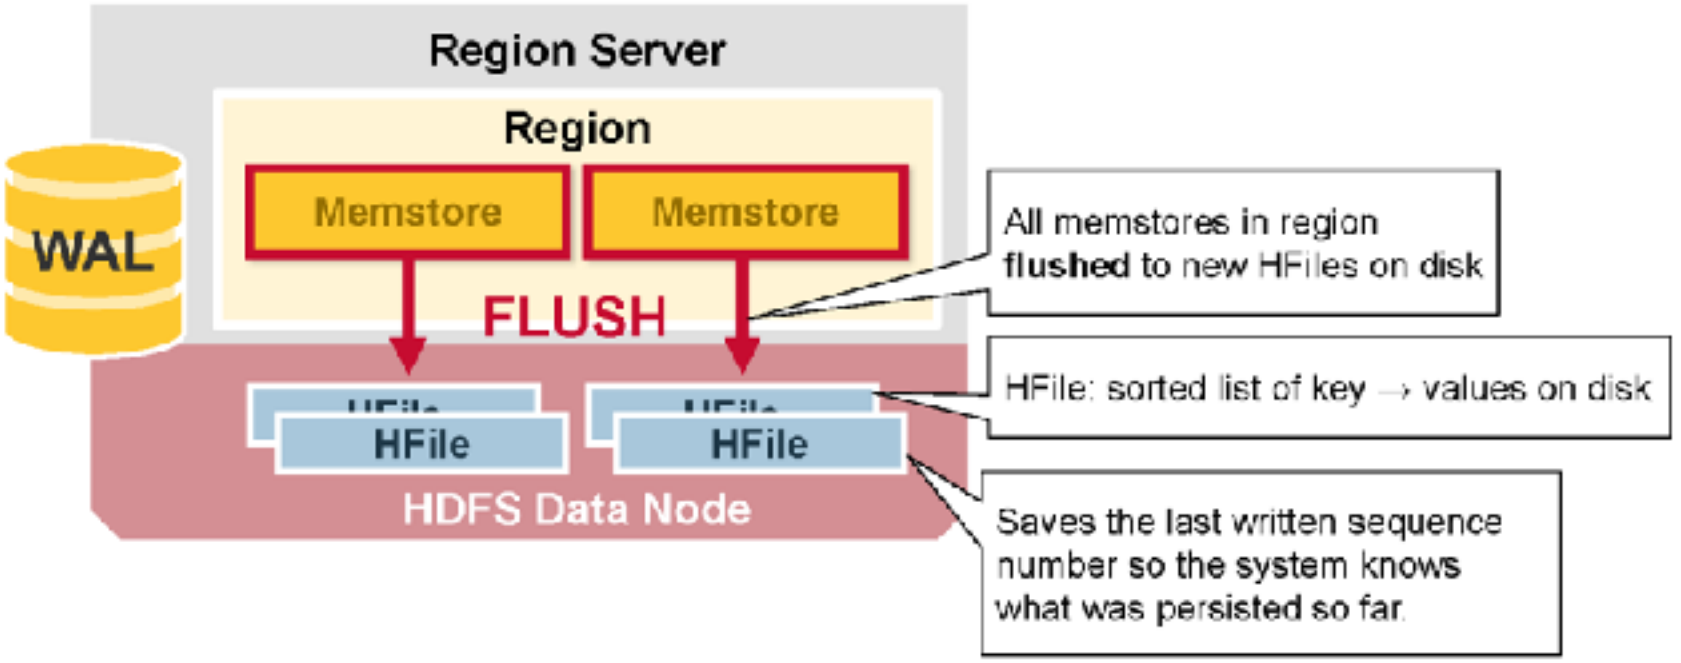

HBase HFile

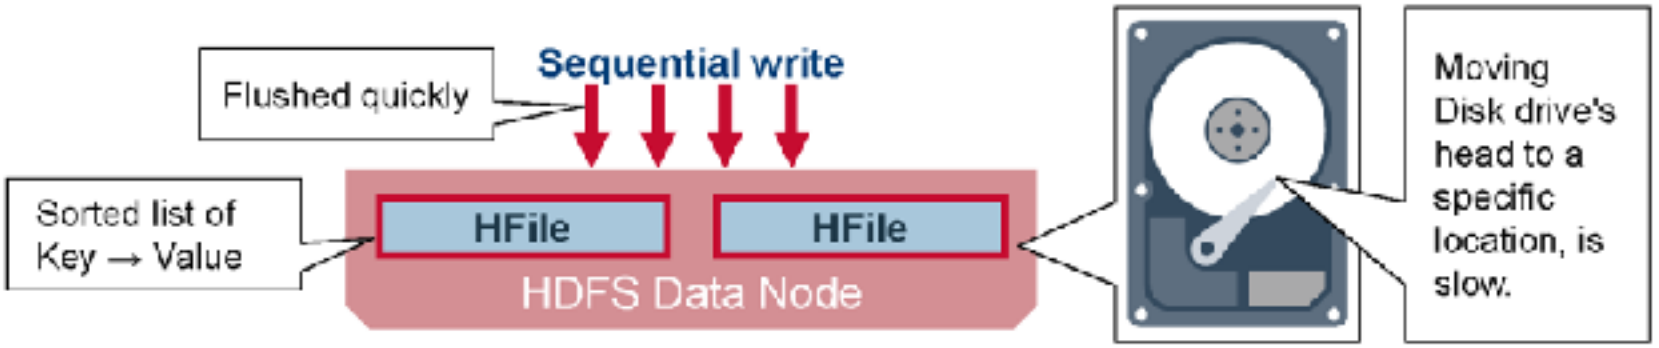

Key Value

Key	CF1:Col	version	value
ra	cf1:ca	v1	1
rb	cf1:cb	v2	4
rb	cf1:cb	v1	3
rc	cf1:ca	v1	5

Key Value

Key	CF2:Col	version	value
ra	cf2:ca	v1	2
rc	cf2:ca	v2	7
rc	cf2:ca	v1	6
rc	cf2:cd	v1	8

HBase HFile Structure

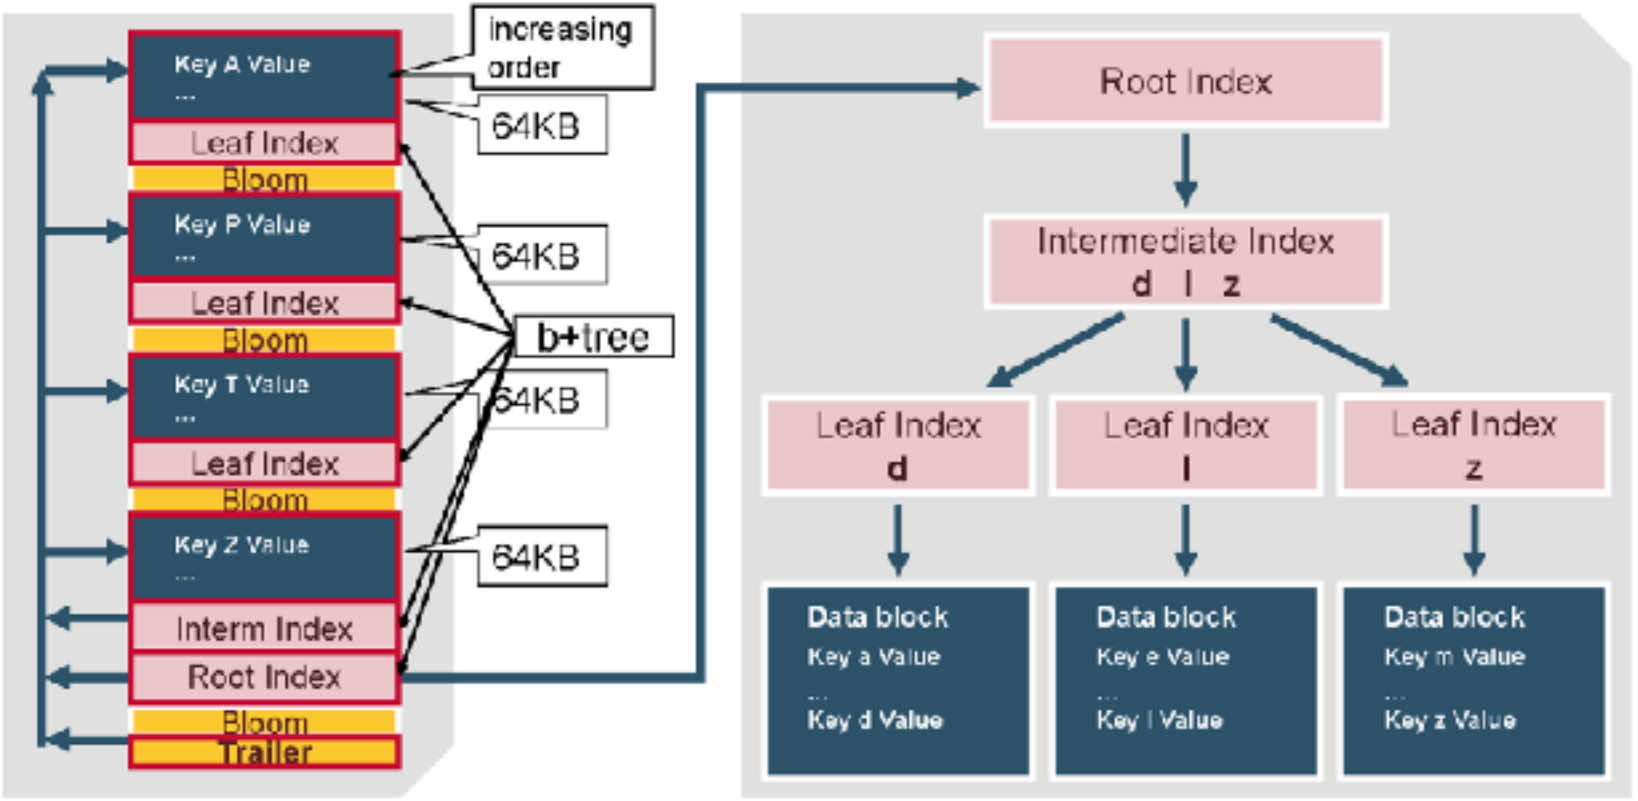

HFile Index

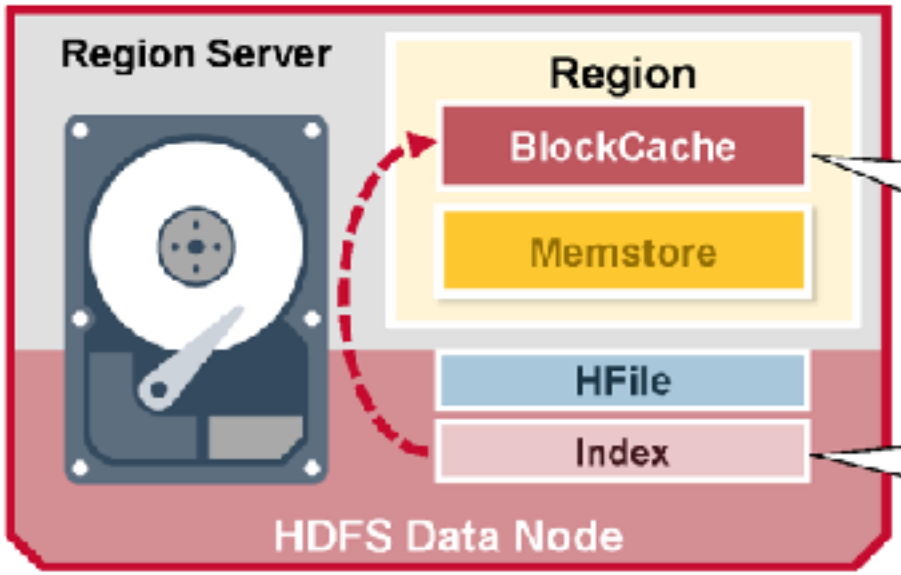

Lookups can be performed with a single disk seek

Hfile Index is kept in Block Cache (Memory)

Index is loaded when the HFile is opened

HBase Read Merge

1. First the scanner looks for the Row KeyValues in the Block cache
2. Next the scanner looks in the Memstore
3. If all row cells not in memstore or blockCache, look in HFiles

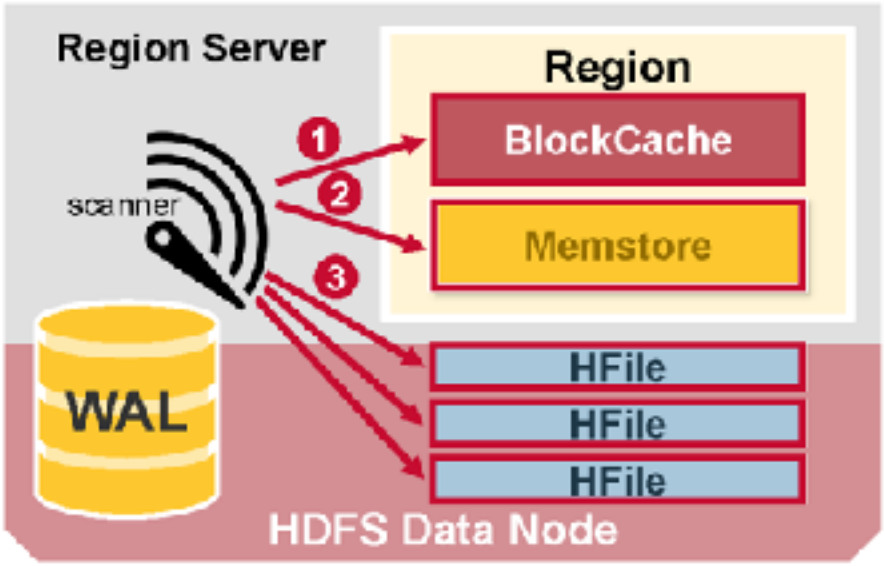

HBase Read Merge

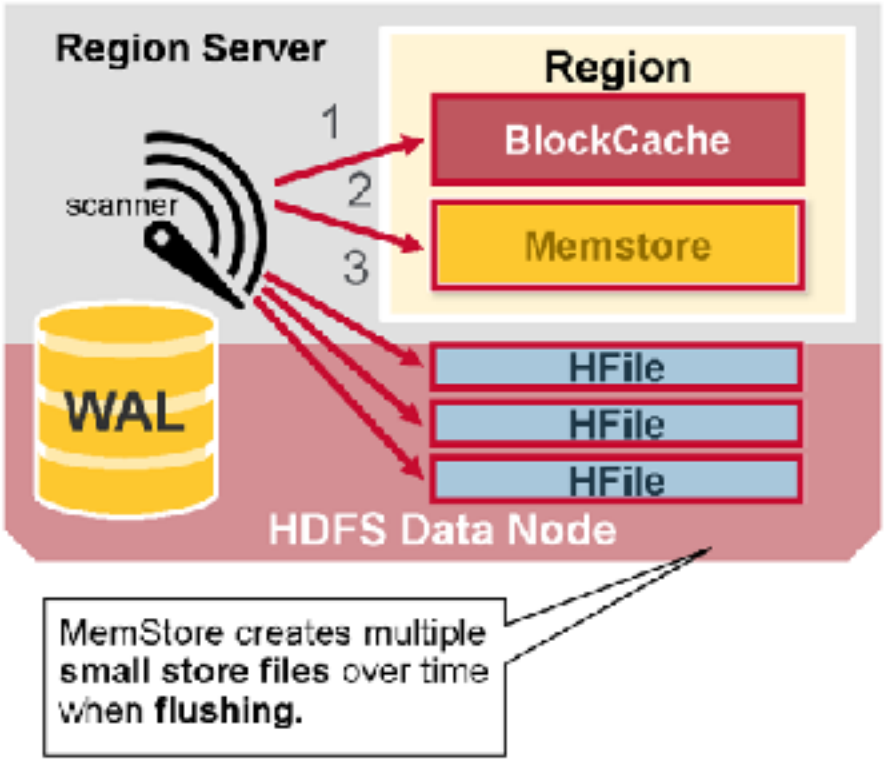

Read Amplification

- multiple files have to be examined

HBase Minor Compaction

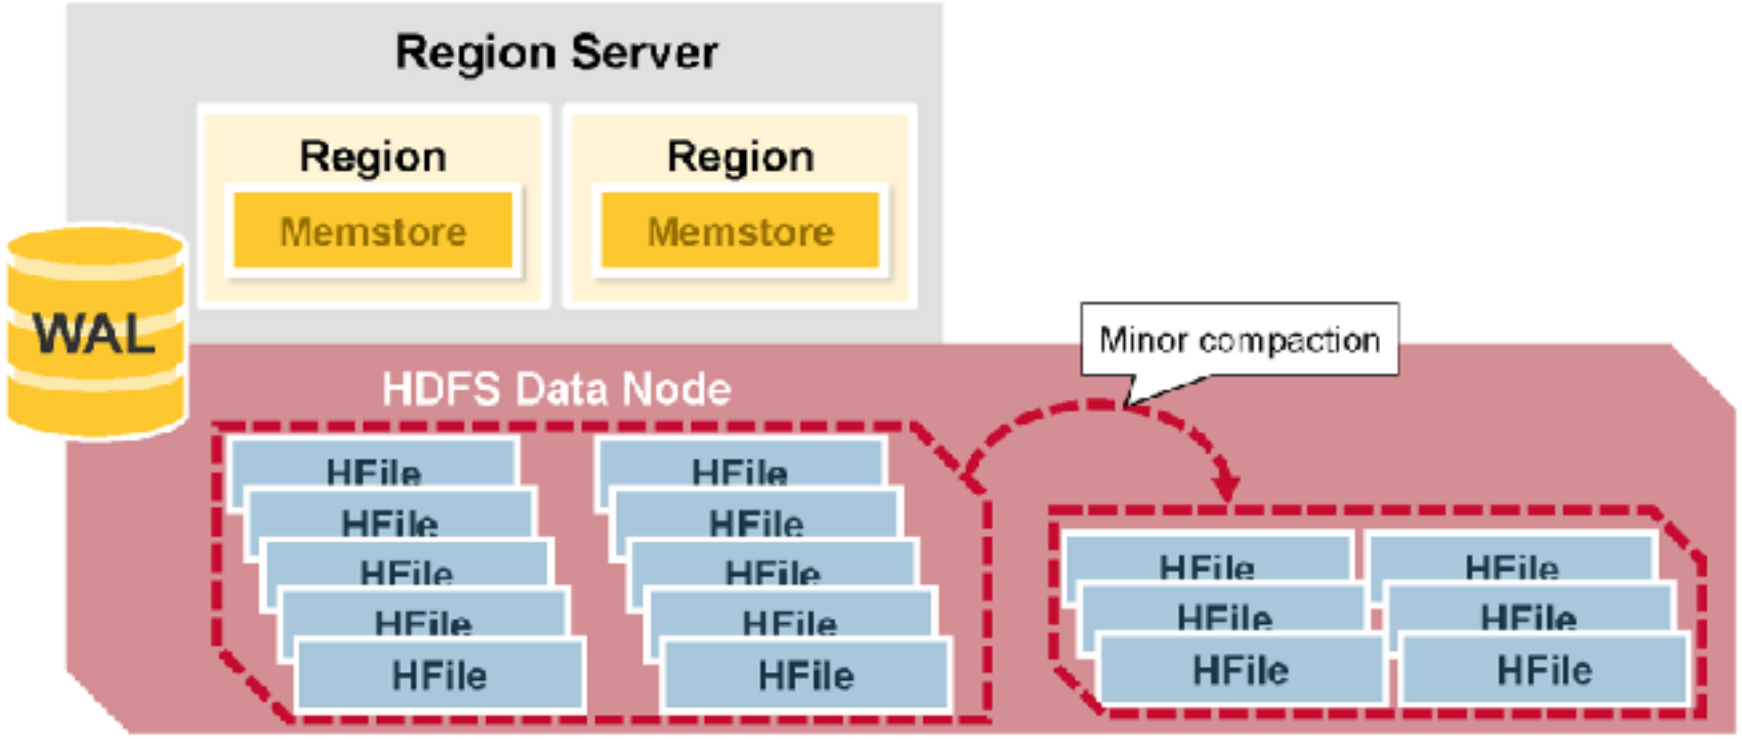

HBase Major Compaction

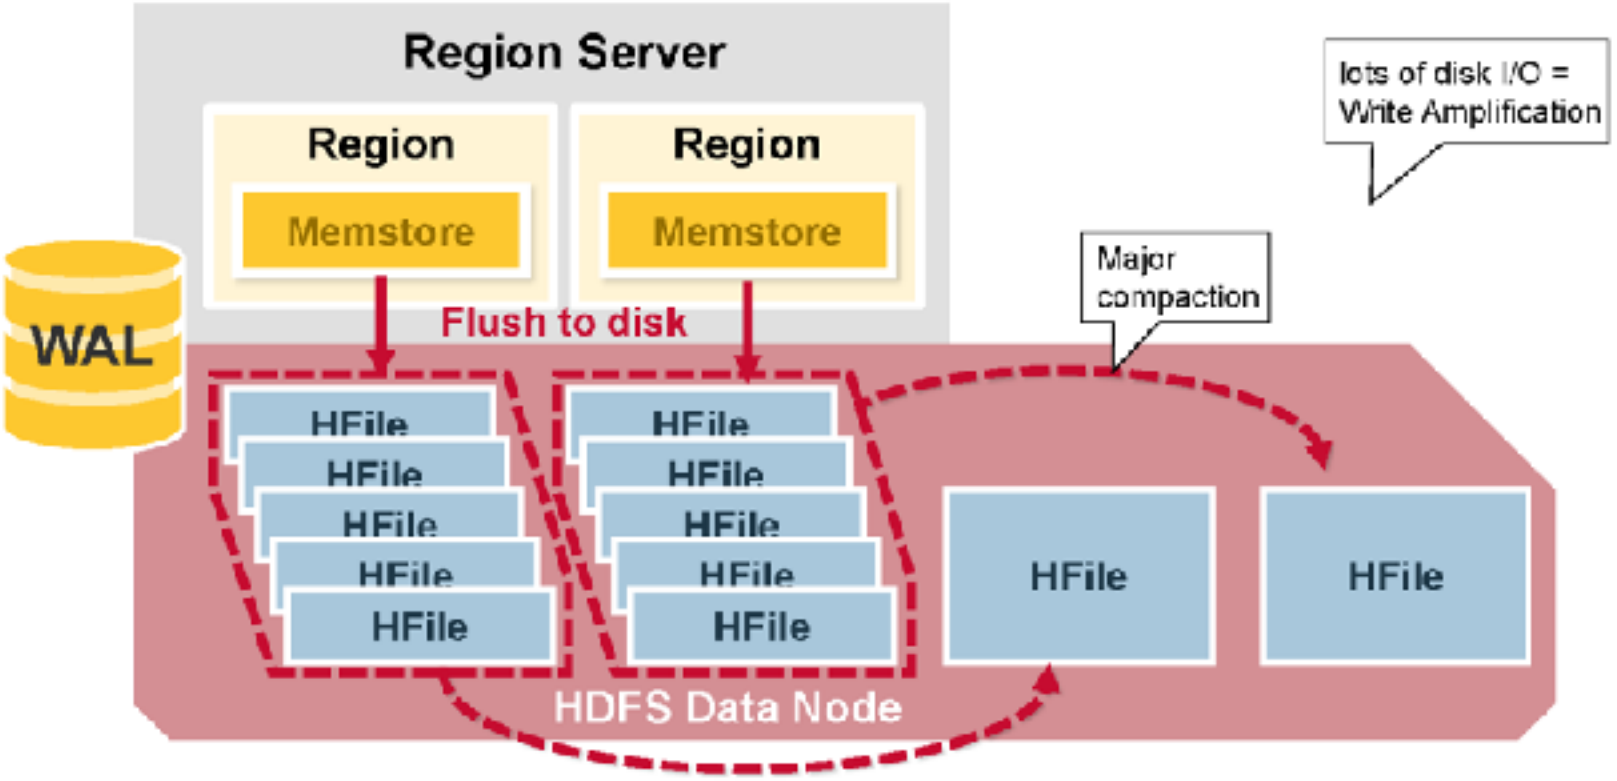

Region = contiguous keys

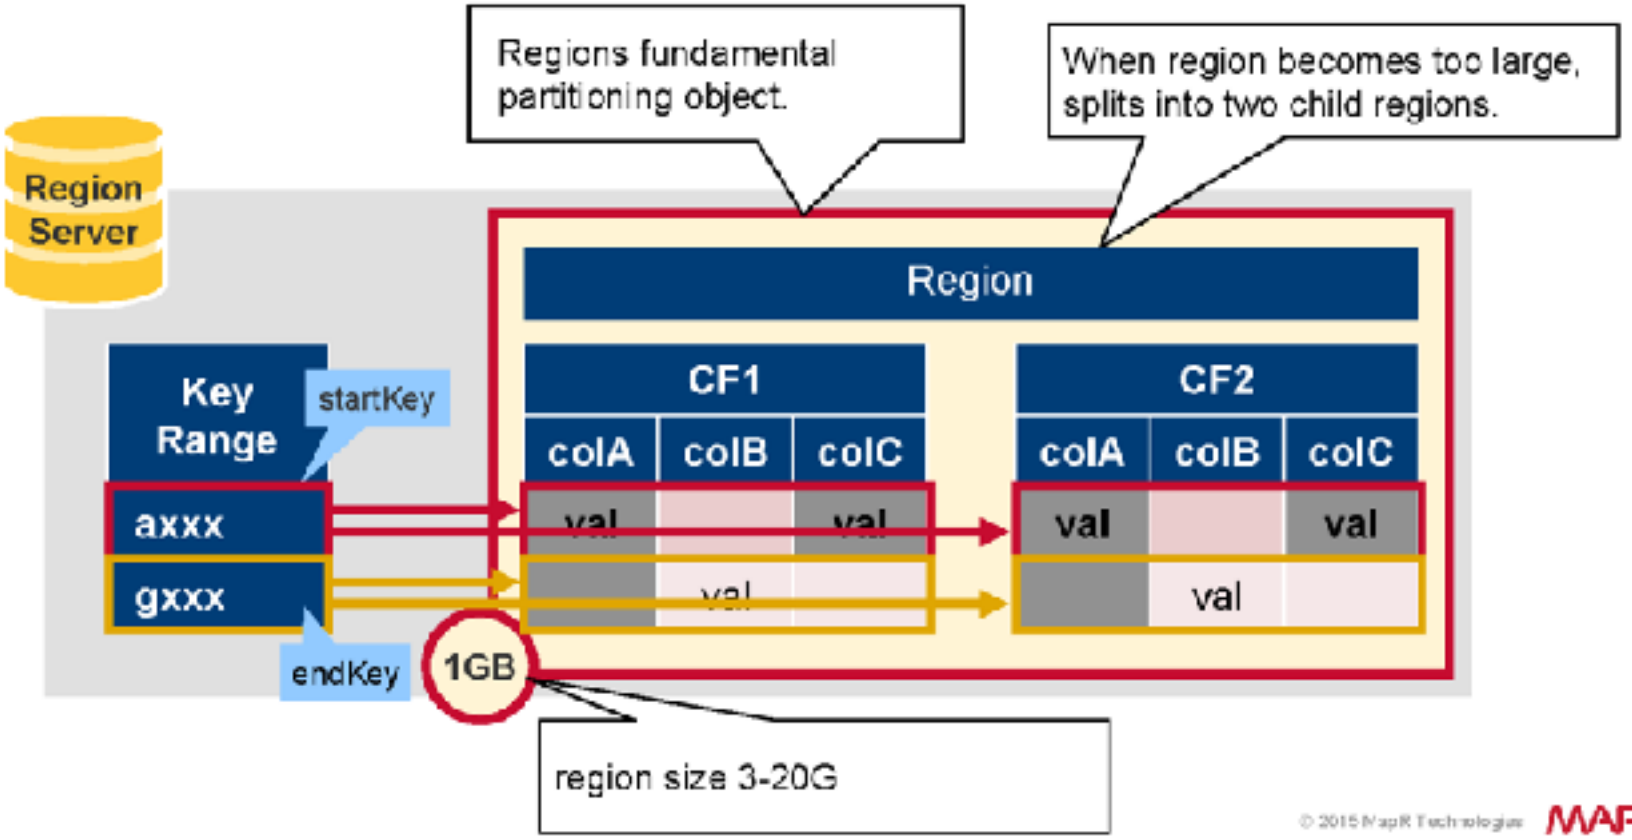

Region Split

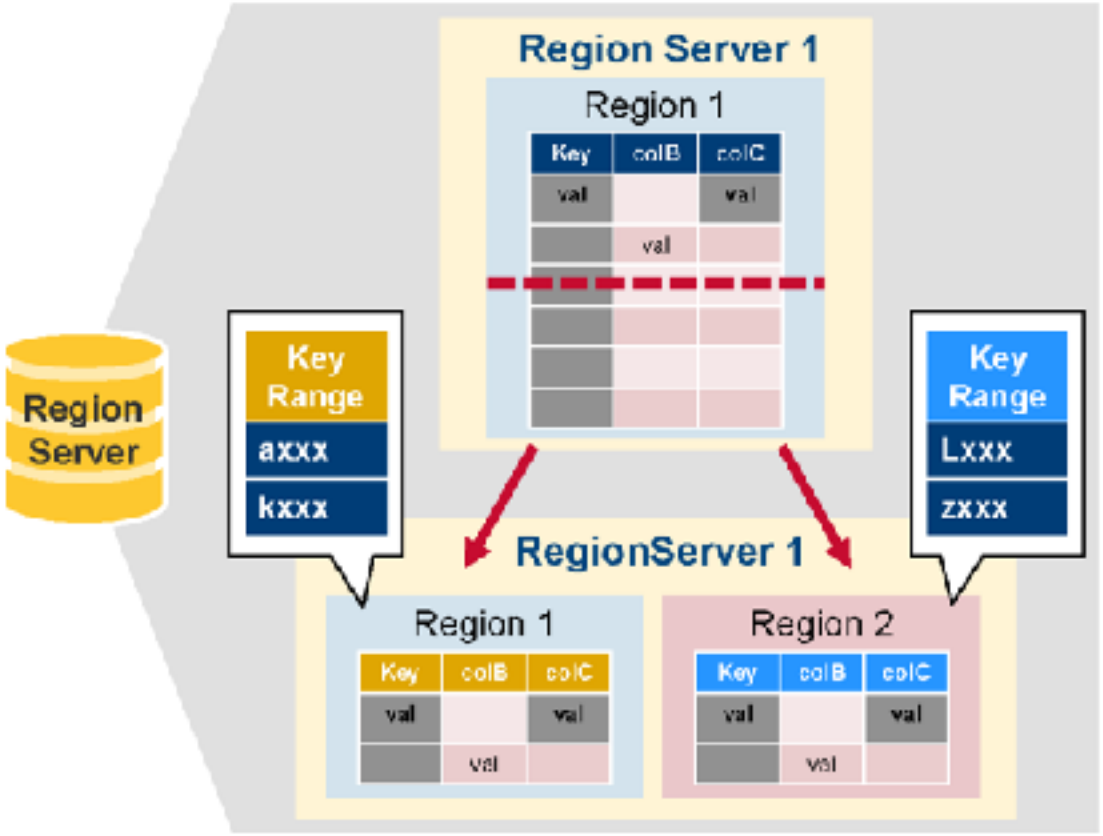

when region size > hbase.hregion.max.filesize → split

Region Load Balancing

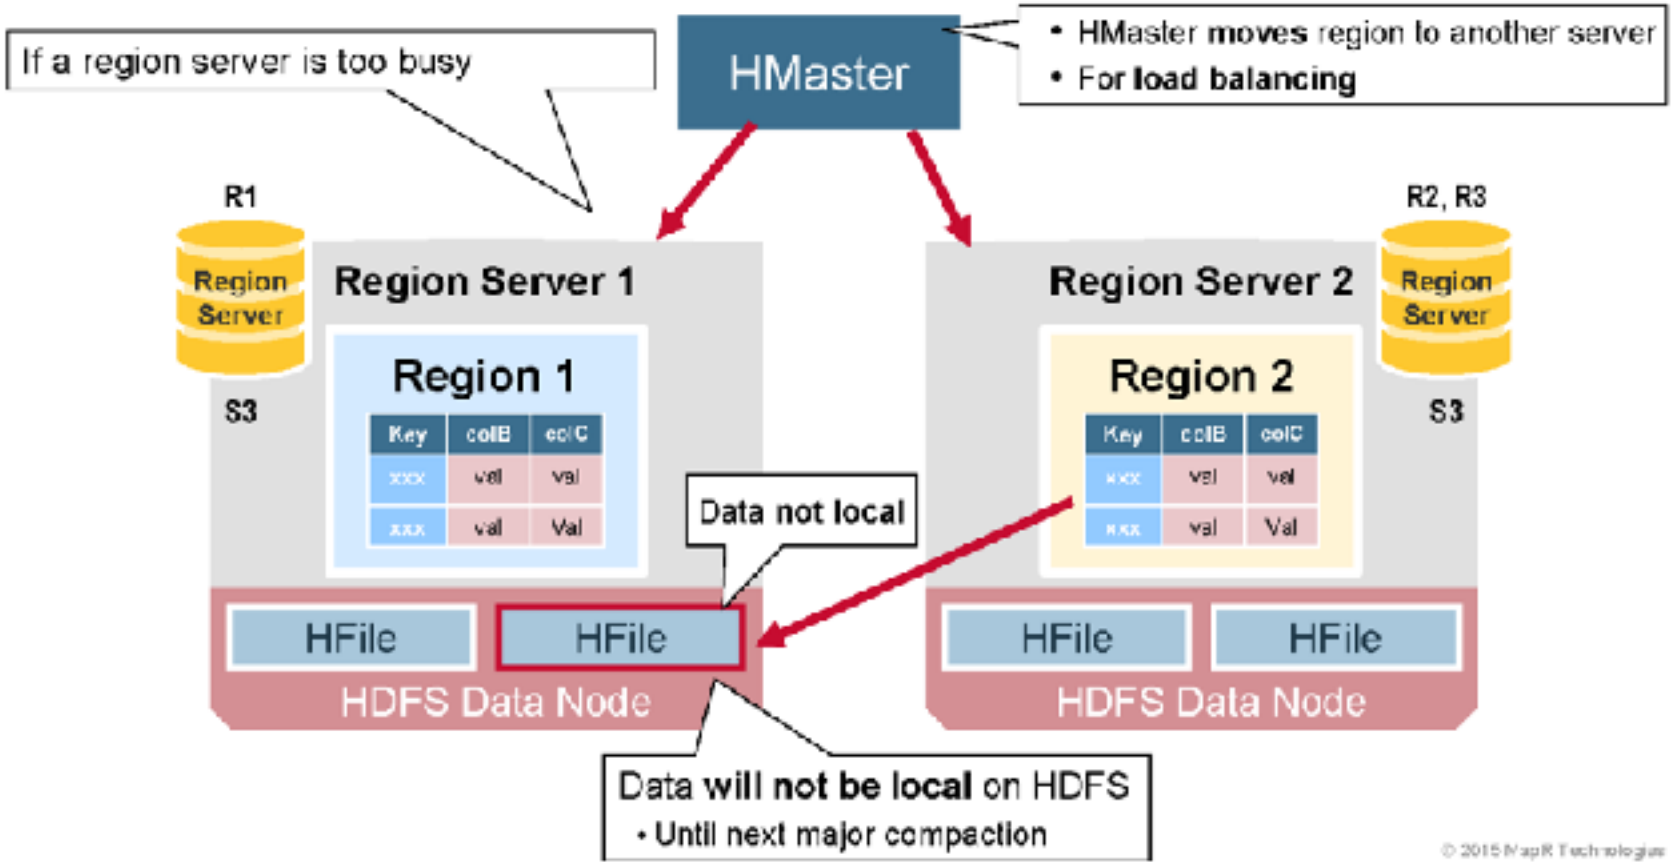

HDFS Data Replication

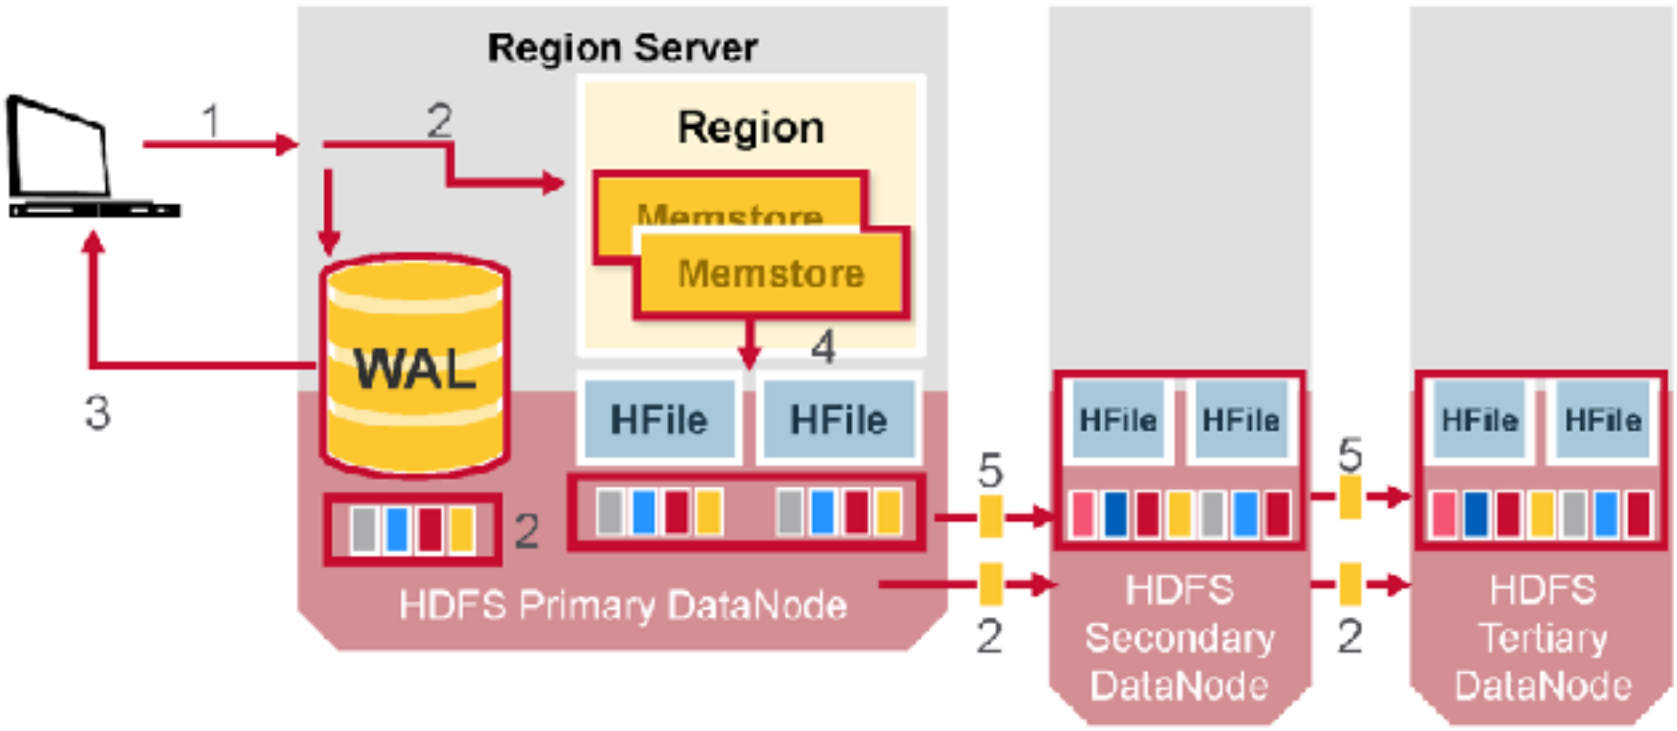

HDFS Data Replication – (2)

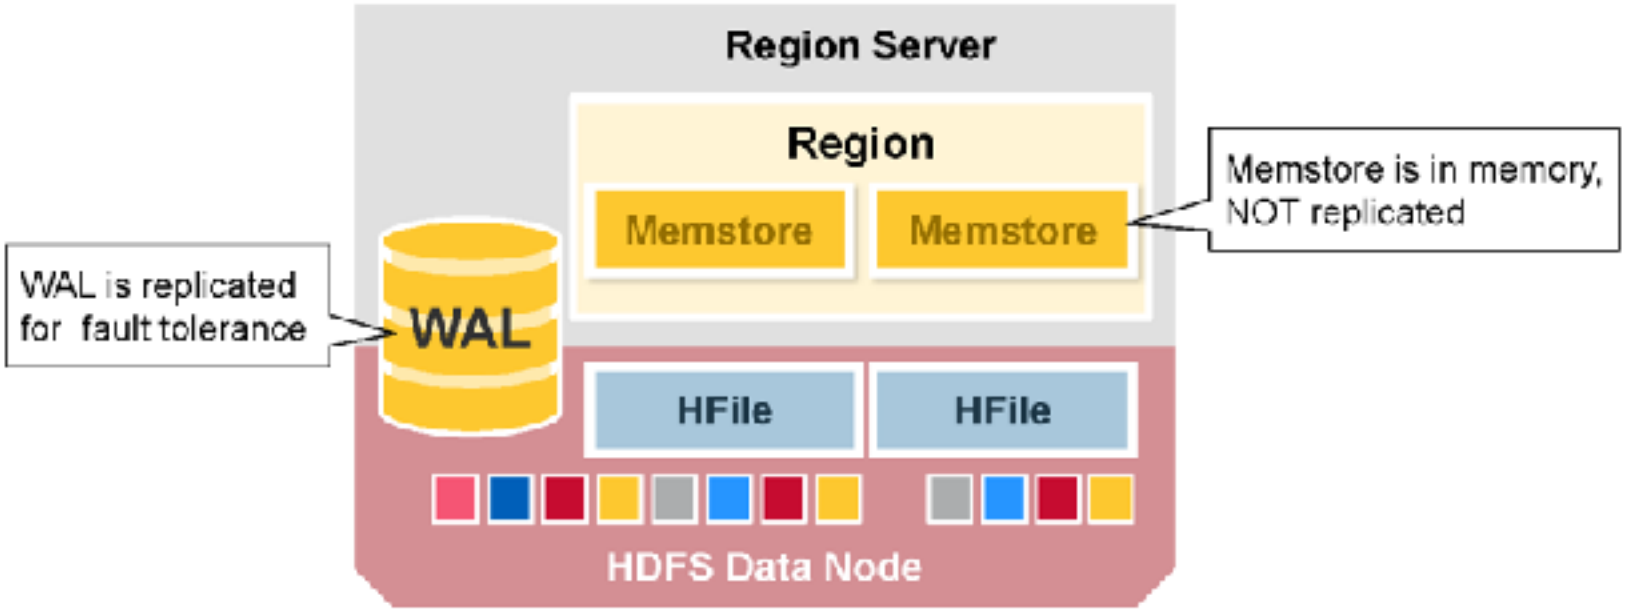

How does HBase recover updates not persisted to HFiles?

HBase Crash Recovery

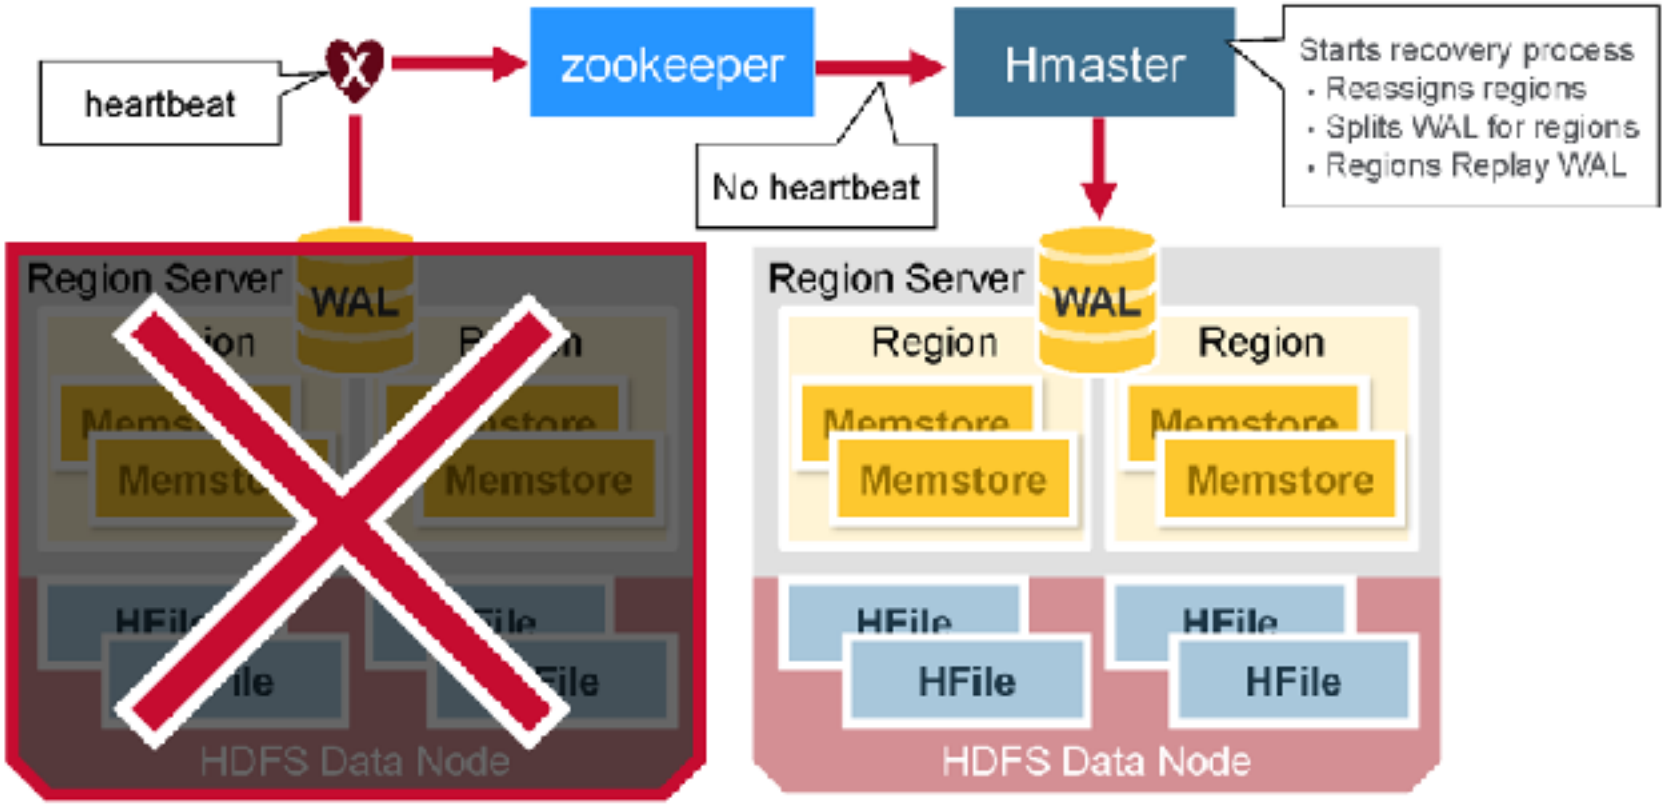

Data Recovery

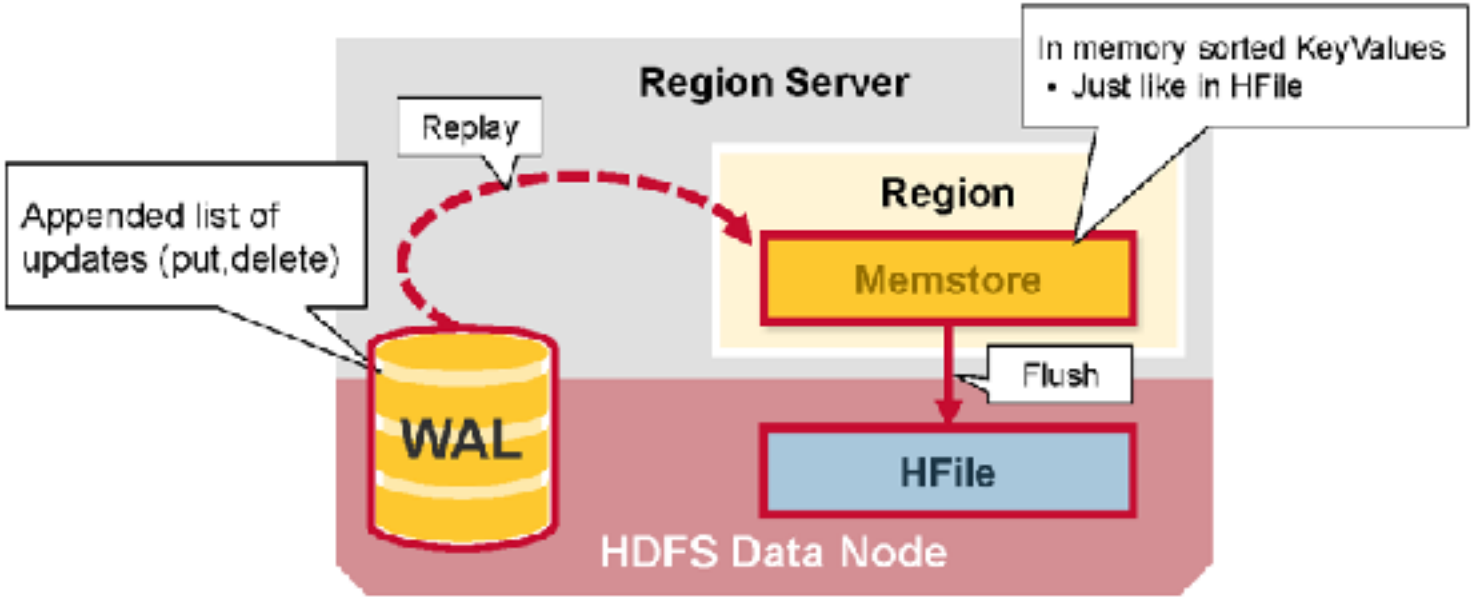

Apache HBase Architecture Benefits

Better than many NoSQL data store solutions, hence its popularity

- **Strong consistency model**
  - When a write returns, all readers will see same value
- **Scales automatically**
  - Regions **split** when data grows too large
  - Uses HDFS to **spread** and **replicate data**
- **Built-in recovery**
  - Using **Write Ahead Log** (similar to journaling on file system)
- **Integrated with Hadoop**
  - MapReduce on HBase is straightforward